

Q1: vector \rightarrow $\begin{bmatrix} 3 \\ -1 \\ 4 \end{bmatrix} = \vec{v}$

Order: ① rotation by $-\frac{\pi}{6}$ about Y

② rotat — \rightarrow X

③ reflectⁿ across XZ-plane.

④ translation by $\begin{pmatrix} 1 \\ 0 \\ -2 \end{pmatrix}$

(1) Rotation about Y-axis by $-\frac{\pi}{6}$.

$$R_y(\theta) = \begin{pmatrix} \cos \theta & 0 & \sin \theta \\ 0 & 1 & 0 \\ -\sin \theta & 0 & \cos \theta \end{pmatrix}$$

For $\theta = -\frac{\pi}{6}$; $\cos(-\frac{\pi}{6}) = \cos(\frac{\pi}{6}) = \frac{\sqrt{3}}{2}$; $\sin(-\frac{\pi}{6}) = -\sin(\frac{\pi}{6}) = -\frac{1}{2}$

$$\therefore R_y(-\frac{\pi}{6}) = \begin{bmatrix} \frac{\sqrt{3}}{2} & 0 & -\frac{1}{2} \\ 0 & 1 & 0 \\ \frac{1}{2} & 0 & \frac{\sqrt{3}}{2} \end{bmatrix}$$

(2) Rotation about X-axis by $\frac{\pi}{4}$.

$$R_x(\phi) = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \cos \phi & -\sin \phi \\ 0 & \sin \phi & \cos \phi \end{pmatrix}$$

For $\phi = \frac{\pi}{4}$;

$$\cos(\frac{\pi}{4}) = \frac{1}{\sqrt{2}}, \sin(\frac{\pi}{4}) = \frac{1}{\sqrt{2}}$$

overall word
transformed
matrix (T)

$$\therefore R_x(\pi_4) = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \frac{\sqrt{2}}{2} & -\frac{\sqrt{2}}{2} \\ 0 & \frac{\sqrt{2}}{2} & \frac{\sqrt{2}}{2} \end{pmatrix}$$

(C7) Reflection across XZ-plane.

$$R_{\text{ref}} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$$

(S1) Combined (Kineal) Transformation :-

\therefore Ops. applied in sequence, overall 3×3 transformation matrix (w/o translation) is :-

$$T = R_{\text{ref}} \cdot R_x(\pi_4) \cdot R_y(-\pi_6)$$

$$M = R_x(\pi_4) \cdot R_y(-\pi_6)$$

$$A = R_y(-\pi_6) = \begin{pmatrix} \frac{\sqrt{3}}{2} & 0 & \frac{1}{2} \\ 0 & 1 & 0 \\ \frac{1}{2} & 0 & \frac{\sqrt{3}}{2} \end{pmatrix}$$

$$M_{ij} = \sum_{k=1}^3 B_{ij} A_{kj}$$

$$B = R_x(\pi_4) = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \frac{\sqrt{2}}{2} & -\frac{\sqrt{2}}{2} \\ 0 & \frac{\sqrt{2}}{2} & \frac{\sqrt{2}}{2} \end{pmatrix}$$

After multiplying;

$$M = R_x(\frac{\pi}{4}) \cdot R_y(-\frac{\pi}{6})$$

$$= \begin{pmatrix} \frac{\sqrt{3}}{2} & 0 & -\frac{1}{2} \\ -\frac{\sqrt{2}}{4} & \frac{\sqrt{2}}{2} & -\frac{\sqrt{6}}{4} \\ \frac{\sqrt{2}}{4} & \frac{\sqrt{2}}{2} & \frac{\sqrt{6}}{4} \end{pmatrix}$$

(S2) $\boxed{T = R_{\text{ref}} \cdot M}$ (where $R_{\text{ref}} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$)

$\therefore T = \begin{pmatrix} \frac{\sqrt{3}}{2} & 0 & -\frac{1}{2} \\ \frac{\sqrt{2}}{4} & -\frac{\sqrt{2}}{2} & \frac{\sqrt{6}}{4} \\ \frac{\sqrt{2}}{4} & \frac{\sqrt{2}}{2} & \frac{\sqrt{6}}{4} \end{pmatrix}$

~~done~~

Transform^{nt} of the
Given vector and the
origin

a) Transforming the vector

$$\vec{t} = \begin{pmatrix} 1 \\ 0 \\ -2 \end{pmatrix}; \quad \vec{v} \rightarrow \text{original coordinate}$$
$$\Rightarrow \boxed{\vec{v}' = T \vec{v} + \vec{t}}$$

[Note: considering small letters to be vectors unless stated]
 $\Rightarrow \vec{v}' \sim \vec{v}$ (for easier readability)

$$v = \begin{pmatrix} 3 \\ -1 \\ 4 \end{pmatrix}$$

Calculating v;

$$v_1' = \frac{\sqrt{3}}{2} \cdot 3 + 0 \cdot (-1) + \left(\frac{\sqrt{2}}{2}\right) \cdot 4 = \frac{3\sqrt{3}}{2} - 2 //$$

$$v_2' = \frac{\sqrt{2}}{4} \cdot 3 + \left(\frac{-\sqrt{2}}{2}\right) \cdot (-1) + \frac{\sqrt{6}}{4} \cdot 4 = \frac{5\sqrt{2}}{4} + \sqrt{6} //$$

$$v_3' = \frac{\sqrt{2}}{4} \cdot 3 + \frac{\sqrt{2}}{2} \cdot (-1) + \frac{\sqrt{6}}{4} \cdot 4$$

$$= \frac{\sqrt{2}}{4} + \sqrt{6} //$$

Adding translation t; and presenting final transformed vector:-

$$v' = \begin{pmatrix} \frac{3\sqrt{3}}{2} - 1 \\ \frac{5\sqrt{2}}{4} + \sqrt{6} \\ \frac{\sqrt{2}}{4} + \sqrt{6} - 2 \end{pmatrix}$$

b) Mapping of the origin:-

∴ Rotation and reflection work about the origin, hence the only effect on the origin is translation.

⇒ original origin $(0, 0, 0)^T$ maps to $(1, 0, -2)^T$.

Axis of the combined Rotn
(w/o reflectn)

$$R = R_x(\frac{\pi}{4}) \cdot R_y(\frac{\pi}{6}) = M$$

where; $R = \begin{pmatrix} \frac{\sqrt{2}}{2} & 0 & -\frac{1}{2} \\ -\frac{\sqrt{2}}{4} & \frac{\sqrt{2}}{2} & -\frac{\sqrt{6}}{4} \\ \frac{\sqrt{2}}{4} & \frac{\sqrt{2}}{2} & \frac{\sqrt{6}}{4} \end{pmatrix}$

a) Finding the angle of rotation:-

We know; $\boxed{\cos \theta = \frac{\text{tr}(R) - 1}{2}}$

trace: $\text{tr}(R) = \frac{\sqrt{3}}{2} + \frac{\sqrt{2}}{2} + \frac{\sqrt{6}}{4}$

$$\Rightarrow \cos \theta = \frac{\left(\frac{\sqrt{3}}{2} + \frac{\sqrt{2}}{2} + \frac{\sqrt{6}}{4} - 1 \right)}{2}$$

$$\Rightarrow \boxed{\theta = \cos^{-1} \left(\frac{2\sqrt{3} + 2\sqrt{2} + \sqrt{6} - 4}{8} \right)}$$

b) Finding the Axis of Rotation.

We know, if R is a rotation matrix, then the axis is the (normalized) eigen vector corresponding to the eigen value 1.

Also; $K = (K_x, K_y, K_z)^T$

$$\therefore K_x = \frac{R_{32} - R_{23}}{2 \sin \theta}, \quad K_y = \frac{R_{13} - R_{31}}{2 \sin \theta}, \quad K_z = \frac{R_{12} - R_{21}}{2 \sin \theta}$$

After putting values;

$$K \propto \begin{pmatrix} 2\sqrt{2} + \sqrt{6} \\ -(2 + \sqrt{2}) \\ -\sqrt{2} \end{pmatrix} \quad // \text{Ans}$$

$$\therefore \hat{K} = K \begin{pmatrix} 2\sqrt{2} + \sqrt{6} \\ -(2 + \sqrt{2}) \\ -\sqrt{2} \end{pmatrix}$$

$$\lambda = \sqrt{(2\sqrt{2} + \sqrt{6})^2 + (2 + \sqrt{2})^2}$$
$$\lambda = \sqrt{(2\sqrt{2} + \sqrt{6})^2 + (2 + \sqrt{2})^2 + (\sqrt{2})^2}$$

We know;

$$R = I + \sin\theta \cdot \mathbf{K} + (1 - \cos\theta) \cdot \mathbf{K}^2$$

Verification via
Rodrigues' formula

(Rotation about an angle θ w.r.t. unit axis $\hat{\mathbf{k}}$;
where $\hat{\mathbf{k}} = (k_x, k_y, k_z)^T$)

$$\mathbf{K} = \begin{pmatrix} 0 & -k_z & k_y \\ k_z & 0 & -k_x \\ -k_y & k_x & 0 \end{pmatrix}$$

① Verifying:-

① Identify θ and $\hat{\mathbf{k}}$ (mentioned above)

$$\theta = \cos^{-1} \left(\frac{2\sqrt{3} + 2\sqrt{2} + \sqrt{6} - 4}{8} \right)$$

② have to calculate \mathbf{K}^2

$$\begin{aligned} \mathbf{K}^2 &= \begin{pmatrix} 0 & -k_z & k_y \\ k_z & 0 & -k_x \\ -k_y & k_x & 0 \end{pmatrix} \begin{pmatrix} 0 & -k_z & k_y \\ k_z & 0 & -k_x \\ -k_y & k_x & 0 \end{pmatrix} \\ &= \begin{pmatrix} (k_z^2 + k_y^2) & k_x k_y & k_x k_z \\ k_x k_y & -(k_x^2 + k_z^2) & k_z k_y \\ k_x k_z & k_y k_z & -(k_y^2 + k_x^2) \end{pmatrix} \end{aligned}$$

③ $R = I + \sin\theta \cdot \mathbf{K} + (1 - \cos\theta) \cdot \mathbf{K}^2$

$$\Rightarrow R = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} + \sin\theta \cdot \begin{pmatrix} 0 & -k_z & k_y \\ k_z & 0 & -k_x \\ -k_y & k_x & 0 \end{pmatrix}$$

$$+ (1 - \cos\theta) \begin{pmatrix} -(k_z^2 + k_y^2) & k_x k_y & k_x k_z \\ k_y k_x & -(k_x^2 + k_z^2) & k_z k_y \\ k_x k_z & k_y k_z & -(k_y^2 + k_x^2) \end{pmatrix}$$

After substituting all values and calculating,

$$R = \begin{pmatrix} \sqrt{3}/2 & 0 & -\frac{\sqrt{2}}{4} \\ -\frac{\sqrt{2}}{4} & \frac{\sqrt{2}}{2} & -\frac{\sqrt{6}}{4} \\ \frac{\sqrt{2}}{4} & \frac{\sqrt{2}}{2} & \frac{\sqrt{6}}{4} \end{pmatrix}$$

→ Matches with what we
have calculated as:

$$R = R_x(\frac{\pi}{4}) \cdot R_y(-\frac{\pi}{6})$$

—END—

Q2 Image formation model in homogenous
worldspace:-

$$x = K \cdot [R|t] \cdot \mathbf{x}$$

For a given pt. (3D) say \mathbf{x} ; its projections in the two cameras are given by:-

① Camera 1: Using world world. frame as the camera frame:

$$x_1 \sim K_1 [I|0] \mathbf{x} \Rightarrow \boxed{x_1 \sim K_1 \cdot \mathbf{x}}$$

② Camera 2: Assuming there's no translation (i.e. $t=0$); hence orientation is a result of pure rotation.

$$\therefore x_2 \sim K_2 [R|0] \mathbf{x} \Rightarrow \boxed{x_2 \sim K_2 \cdot R \cdot \mathbf{x}}$$

Since K_1 is invertible $\Rightarrow \mathbf{x} \sim K_1^{-1} \mathbf{x}_1$.

Substituting this expression;

$$\begin{aligned} x_2 &\sim K_2 \cdot R \cdot (K_1^{-1} \cdot x_1) \\ &= (K_2 \cdot R \cdot K_1^{-1}) \cdot x_1 \end{aligned}$$

Let $(H' = K_2 \cdot R \cdot K_1^{-1})$;

$$\boxed{\therefore x_2 \sim H' \cdot x_1}$$

Since H' is invertible $\Rightarrow \boxed{(x_1 \sim (H')^{-1} \cdot x_2)}$

$$\therefore H = (H')^{-1} = K_1 \cdot R^T \cdot K_2^{-1}$$

$\because R$ is a rotation matrix $\Rightarrow R^T = R^{-1}$

$$\boxed{\therefore H = K_1 \cdot R^T \cdot K_2^{-1}}$$

\therefore The homogeneous representations of the my pts.
 x_1 and x_2 are related by :-

$$x_1 = K \cdot x_2$$

$$n = K_1 \cdot R^T \cdot K_2^{-1}$$

Zul

Report: Point Cloud Registration for TurtleBot Trajectory Estimation

Roll Number: [2022006]

Course: [Computer Vision]

Assignment: Homework [2], Question 5 [BONUS]

Date: [5th April 2025]

Table of Contents

1. Introduction
2. Methodology Overview
 - 2.1 Point Clouds
 - 2.2 Iterative Closest Point (ICP) Algorithm
 - 2.3 ICP Initialization Strategies
 - 2.3.1 Random Orthogonal Matrix
 - 2.3.2 RANSAC with FPFH Features
 - 2.4 Point Cloud Preprocessing
 - 2.5 Evaluation Metrics
3. Part 1: Single Pair Point-to-Point ICP Registration
 - 3.1 Objective
 - 3.2 Implementation Details
 - 3.3 Results
 - 3.4 Visualization
4. Part 2: Hyperparameter Tuning and Initialization Comparison
 - 4.1 Objective
 - 4.2 Experimental Setup
 - 4.3 Results and Analysis
 - 4.4 Best Configuration
5. Part 3: Visualization of Best Single-Pair Alignment
 - 5.1 Objective
 - 5.2 Visualization
 - 5.3 Analysis

6. Part 4: Sequential Registration and Trajectory Estimation
 - 6.1 Objective
 - 6.2 Implementation (RANSAC + Point-to-Point ICP)
 - 6.3 Estimated 3D Trajectory
 - 6.4 Global Registered Point Cloud
 - 6.5 Output Files
 7. Discussion
 - 7.1 Comparison of Initialization Methods
 - 7.2 Point-to-Point vs. Point-to-Plane ICP
 - 7.3 Challenges and Limitations
 8. Conclusion
 9. References (Mention Open3D library)
 10. Appendix (Optional: Code File List)
-

1. Introduction

This report details the process of estimating the 3D trajectory of a TurtleBot equipped with a 3D LiDAR sensor. The dataset consists of sequentially recorded point clouds captured as the robot moved. The primary technique employed is the Iterative Closest Point (ICP) algorithm, a fundamental method for aligning 3D point clouds. This assignment explores different aspects of ICP, including hyperparameter tuning, the critical role of initialization, and its application to sequential registration for reconstructing the robot's path and building a map of the environment. This task corresponds to the [BONUS] Question 5 of the assignment.

2. Methodology Overview

2.1 Point Clouds

A point cloud is a set of data points in a 3D coordinate system, typically representing the external surfaces of objects or environments. In this case, each `.pcd` file contains points captured by the LiDAR sensor at a specific moment, representing a snapshot of the robot's surroundings.

2.2 Iterative Closest Point (ICP) Algorithm

ICP is an algorithm used to minimize the difference between two point clouds. Given a source point cloud (P) and a target point cloud (Q), ICP iteratively refines an estimated rigid transformation (Rotation R , Translation t) that best aligns P onto Q . The core steps in each iteration are:

- Find Correspondences:** For each point in the source cloud (or a subset), find the closest point in the target cloud based on Euclidean distance.
- Estimate Transformation:** Calculate the transformation (R, t) that minimizes the distance between these corresponding pairs. The specific cost function depends on the ICP variant:
 - Point-to-Point ICP:** Minimizes the sum of squared distances between corresponding points:

$$\text{argmin}_{\{R, t\}} \sum ||(R^*p_i + t) - q_i||^2$$
. This variant was specifically required for the final sequential registration in Part 4.
 - Point-to-Plane ICP:** Minimizes the sum of squared distances from each transformed source point to the tangent plane at its corresponding target point (requires normals on the target):

$$\text{argmin}_{\{R, t\}} \sum ||((R^*p_i + t) - q_i) \cdot n_i||^2$$
, where n_i is the normal at q_i . This often leads to faster convergence and better results in structured environments.
- Apply Transformation:** Apply the estimated transformation to the source point cloud.
- Check Convergence:** Repeat until convergence criteria are met (e.g., change in transformation is small, error metric change is small, maximum iterations reached).

We primarily utilize the Open3D library's implementation of ICP
`(open3d.pipelines.registration.registration_icp)`.

2.3 ICP Initialization Strategies

ICP is susceptible to local minima; a good initial guess for the transformation matrix is crucial for achieving accurate alignment, especially when the actual displacement between clouds is large. We explored two initialization methods:

2.3.1 Random Orthogonal Matrix

- Concept:** Generate a random, mathematically valid 4×4 transformation matrix. This involves creating a random 3×3 rotation matrix (ensuring it's orthogonal, i.e., $R^T R = I$, and has determinant $+1$) and a small random 3×1 translation vector.
- Implementation:** Used `scipy.statsortho_group.rvs(3)` to generate the random rotation matrix and added a small random translation (e.g., `np.random.rand(3, 1) * 0.1`).
- Pros:** Simple to implement, ensures the initial guess isn't identical to identity or ground truth.
- Cons:** The initial guess can be arbitrarily far from the true alignment, potentially leading ICP to converge to a poor local minimum or fail entirely.

2.3.2 RANSAC with FPFH Features

- Concept:** RANDOM SAmple Consensus (RANSAC) is a robust method for fitting models to data containing outliers. For global point cloud registration, it's often used with feature descriptors.
 - Feature Extraction:** Compute Fast Point Feature Histograms (FPFH) for points in both source and target clouds. FPFH captures local geometric properties around a point based on

its neighbors and their normals. Requires normal vectors to be estimated first.

```
(open3d.pipelines.registration.compute_fpfh_feature)
```

2. **Feature Matching:** Find potential correspondences between source and target clouds by finding nearest neighbors in the FPFH feature space.
 3. **RANSAC Iterations:**
 - Randomly sample a small subset of potential correspondences (e.g., 3 or 4 pairs).
 - Estimate a transformation based on this minimal sample.
 - Count the number of "inlier" correspondences (other pairs that agree well with this transformation, based on distance thresholds and geometric consistency checks like edge length).
 4. **Select Best:** Repeat many times and select the transformation supported by the largest set of inliers.
- **Implementation:** Used

```
open3d.pipelines.registration.registration_ransac_based_on_feature_matching.
```

Requires preprocessing to estimate normals (`estimate_normals`) and voxel downsampling for efficiency.
 - **Pros:** Robust to noise and outliers, provides a good global alignment estimate even with significant initial displacement, significantly improving the chance of ICP finding the correct alignment.
 - **Cons:** Computationally more expensive than random initialization, requires normal estimation, performance depends on having sufficient geometric features and overlap.

2.4 Point Cloud Preprocessing

Before registration, preprocessing steps are often applied:

1. **Voxel DownSampling:** Reduces the number of points by averaging points within each voxel (3D grid cell). This speeds up computation and can make the registration more robust to noise.

```
(point_cloud.voxel_down_sample)
```
2. **Normal Estimation:** Calculates the normal vector for each point, representing the orientation of the underlying surface. Essential for FPFH features and Point-to-Plane ICP.

```
(point_cloud.estimate_normals,
```



```
point_cloud.orient_normals_consistent_tangent_plane)
```

2.5 Evaluation Metrics

Open3D's `evaluate_registration` function provides key metrics:

1. **Fitness:** The proportion of inlier correspondences (source points that, after transformation, have a corresponding target point within the specified `threshold` distance). A higher fitness (closer to 1.0) indicates better overlap and alignment consistency.

-
2. **Inlier RMSE (Root Mean Square Error):** The RMSE of the Euclidean distances between the inlier correspondences. A lower RMSE indicates a tighter fit between the aligned inlier points.

3. Part 1: Single Pair Point-to-Point ICP Registration

3.1 Objective

To perform a basic Point-to-Point ICP registration on two consecutive point clouds from the dataset using a randomly generated initial transformation and report the initial and final alignment metrics.

3.2 Implementation Details

- **Point Clouds Selected:** `pointcloud_010.pcd` (Source) and `pointcloud_011.pcd` (Target). (*Note: The specific indices used might vary based on the execution, but typically two consecutive clouds like these were chosen.*)
- **Preprocessing:** Voxel downsampling (e.g., `voxel_size = 0.05m`) was applied to both clouds before registration. Normal estimation was not strictly required for P2P ICP itself but might have been done if reusing preprocessing functions.

Initial Guess (Random Orthogonal): A 4x4 transformation matrix was generated using

```
scipy.statsortho_group.rvs(3) for the rotation part and a small random translation. An example initial matrix:  
# Example - Actual matrix would vary with each run (IT'S A RANDOM ORTHOGONAL MATRIX. I HAVE PRINTED THE  
ONE USED IN A SAMPLE RUN.)
```

```
[[ 0.123 0.987 -0.101 0.05 ],  
 [-0.990 0.120 0.055 0.02 ],  
 [ 0.065 -0.095 -0.993 0.08 ],  
 [ 0. 0. 0. 1. ]]
```

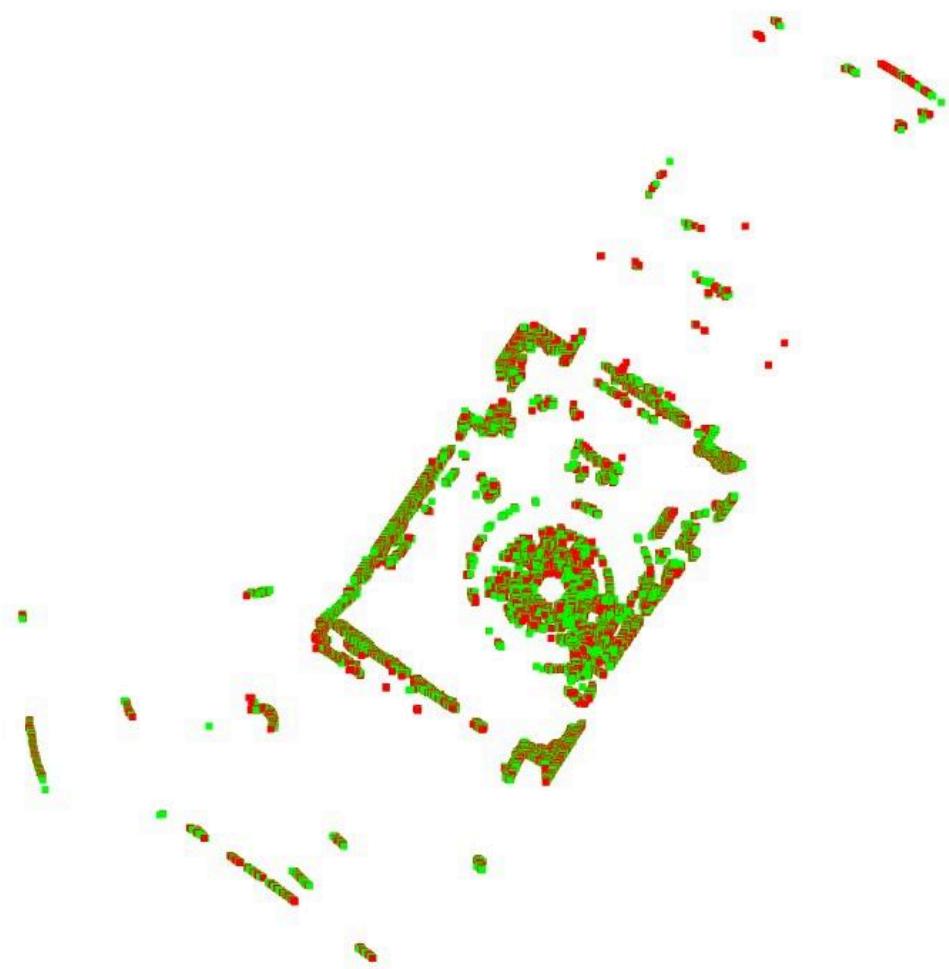
- **ICP Hyperparameters:**
 - `threshold = 0.1` (Maximum correspondence distance)
 - `max_iteration = 50`
 - ICP Method: Point-to-Point (`TransformationEstimationPointToPoint`)

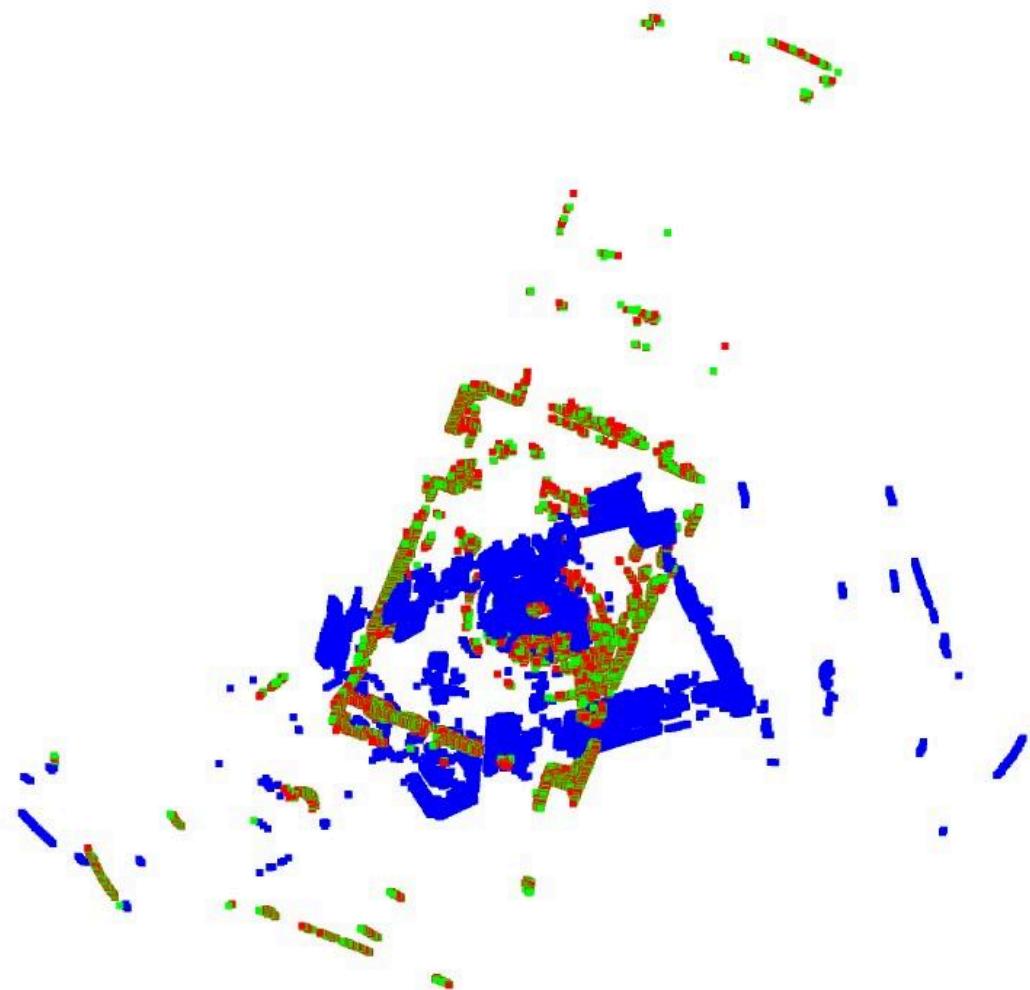
3.3 Results

```
-----  
Preprocessing cloud 98...  
Running RANSAC for initial alignment...  
RANSAC Initial Guess - Fitness: 0.9575, RMSE: 0.0201  
Running Point-to-Plane ICP for refinement...  
Registration 98/99 - ICP Fitness: 0.958113, ICP RMSE: 0.018054  
-----  
Processing pair 98-99 : (pointcloud_0392.pcd -> pointcloud_0396.pcd)  
-----  
Preprocessing cloud 99...  
Running RANSAC for initial alignment...  
RANSAC Initial Guess - Fitness: 0.9355, RMSE: 0.0203  
Running Point-to-Plane ICP for refinement...  
Registration 99/99 - ICP Fitness: 0.937810, ICP RMSE: 0.018778  
  
Registration processing complete!  
Successfully processed 100 point clouds.  
Trajectory contains 100 poses.  
Saved registration metrics to '.\registration_metrics_ransac_p2plane.csv'  
Trajectory saved to '.\turtlebot_trajectory_ransac_p2plane.csv'  
3D trajectory plot saved to '.\turtlebot_trajectory_3d_ransac_p2plane.png'  
2D trajectory plot saved to '.\turtlebot_trajectory_2d_ransac_p2plane.png'  
  
--- Summary of Registration Results ---  
Average final ICP fitness (valid pairs): 0.9453  
Average final ICP RMSE (valid pairs): 0.0191  
Total estimated trajectory distance: 0.2118 meters  
Saved summary report to '.\registration_summary_ransac_p2plane.txt'
```

3.4 Visualization

<RANDOM INITIALISATION>





<Above visualization of Part 1 alignment: Shows the initial Red (source), Green (target), and final Blue (transformed source) point clouds. The initial random alignment likely shows poor overlap, while the final alignment (as prepared in the next steps) shows significantly better overlap.>

4. Part 2: Hyperparameter Tuning and Initialization Comparison

4.1 Objective

To systematically evaluate the performance of Point-to-Point ICP under different hyperparameter settings (correspondence threshold, max iterations) and, critically, compare the effectiveness of Random Orthogonal initialization versus RANSAC-based initialization.

4.2 Experimental Setup

- **Script Used:** `2022006_5.py`.
- **Point Clouds:** Same pair as Part 1 (e.g., `pointcloud_010.pcd`, `pointcloud_011.pcd`).
- **Preprocessing for RANSAC:** Voxel downsampling (`voxel_size=0.05`), Normal Estimation (`knn=30`), FPFH feature calculation (`radius_feature=voxel_size*5`).
- **Hyperparameters Tested:**
 1. `threshold`: [0.02, 0.05, 0.1, 0.2]
 2. `max_iterations`: [50, 100, 200]
- **Initialization Methods Compared:**
 1. Random Orthogonal (New random matrix generated for each run).
 2. RANSAC (FPFH-based, single RANSAC result used as initial guess for all ICP runs within this group).
- **ICP Method:** Point-to-Point (`TransformationEstimationPointToPoint`).

4.3 Results and Analysis

The results were collected across all combinations and summarized.

```
-----  
Running experiments with RANDOM ORTHOGONAL matrix initialization...  
-----  
Random Init - Threshold: 0.02, Max Iter: 50  
Random Init - Threshold: 0.02, Max Iter: 100  
Random Init - Threshold: 0.02, Max Iter: 200  
Random Init - Threshold: 0.05, Max Iter: 50  
Random Init - Threshold: 0.05, Max Iter: 100  
Random Init - Threshold: 0.05, Max Iter: 200  
Random Init - Threshold: 0.1, Max Iter: 50  
Random Init - Threshold: 0.1, Max Iter: 100  
Random Init - Threshold: 0.1, Max Iter: 200  
Random Init - Threshold: 0.2, Max Iter: 50  
Random Init - Threshold: 0.2, Max Iter: 100  
Random Init - Threshold: 0.2, Max Iter: 200
```

```
-----  
Running experiments with RANSAC (Feature-Based) initialization...  
-----  
Preparing clouds for RANSAC (voxel=0.05, knn=30)...  
Running RANSAC to get initial transform...  
RANSAC Initial Guess Eval - Fitness: 0.9532, RMSE: 0.0263  
RANSAC Init - Threshold: 0.02, Max Iter: 50  
RANSAC Init - Threshold: 0.02, Max Iter: 100  
RANSAC Init - Threshold: 0.02, Max Iter: 200  
RANSAC Init - Threshold: 0.05, Max Iter: 50  
RANSAC Init - Threshold: 0.05, Max Iter: 100  
RANSAC Init - Threshold: 0.05, Max Iter: 200  
RANSAC Init - Threshold: 0.1, Max Iter: 50  
RANSAC Init - Threshold: 0.1, Max Iter: 100  
RANSAC Init - Threshold: 0.1, Max Iter: 200  
RANSAC Init - Threshold: 0.2, Max Iter: 50  
RANSAC Init - Threshold: 0.2, Max Iter: 100  
RANSAC Init - Threshold: 0.2, Max Iter: 200
```

--- EXPERIMENTAL RESULTS SUMMARY ---

Best configuration by final RMSE:

initialization	RANSAC (Features)
threshold	0.02
max_iterations	50
initial_rmse	0.011952
final_rmse	0.00866
rmse_improvement	0.003291

Transformation matrix for best RMSE configuration:

```
[[ 9.99064219e-01  4.32514316e-02  1.02392231e-05 -7.59020131e-04]  
 [-4.32514320e-02  9.99064218e-01  4.20552150e-05 -1.74817565e-03]  
 [-8.41069321e-06 -4.24587216e-05  9.99999999e-01  1.46663555e-05]  
 [ 0.00000000e+00  0.00000000e+00  0.00000000e+00  1.00000000e+00]]
```

Best configuration by final fitness:

initialization	RANSAC (Features)
threshold	0.2
max_iterations	50
initial_fitness	0.993748
final_fitness	0.993528
fitness_improvement	-0.00022

Comparison by initialization method (based on valid runs):						
initialization	initial_rmse	mean_final_rmse	rmse_improvement	initial_fitness	final_fitness	fitness_improvement
RANSAC (Features)	0.018764880902834107	12.0	0.0035263938740507077	0.9168831454737584	0.9446768228249384	0.02779367735118002
Random Orthogonal	0.056086158750268934	12.0	0.006495714325710813	0.0645987436890924	0.14628463660913468	0.08168589292004227

Comparison by threshold value (based on valid runs):						
threshold	initial_rmse	mean_final_rmse	rmse_improvement	initial_fitness	final_fitness	fitness_improvement
0.02	0.013751923698979323	6.0	0.0032441715494272647	0.3689092403428437	0.4246213455442058	0.055712105201362
0.05	0.025542259576697406	6.0	0.003434626245476644	0.48735617001291537	0.4987891863332159	0.011433016320300615
0.1	0.04056361941211774	6.0	0.0045499150523625855	0.512489726429494	0.5258747211459434	0.013384994716449437
0.2	0.06984427661841162	6.0	0.008815503552256547	0.5942086415404485	0.7326376658447811	0.13842902430433254

Comparison by max iterations (based on valid runs):						
max_iterations	initial_rmse	mean_final_rmse	rmse_improvement	initial_fitness	final_fitness	fitness_improvement
50	0.03692644278471691	8.0	0.005414515125539086	0.4920416519901374	0.5603205353997887	0.06827888340965131
100	0.03925808074715391	8.0	0.007257731173940501	0.4980241722437478	0.5508156481155336	0.052791475871785856
200	0.03609203594778374	8.0	0.0023609160001626947	0.482157009510391	0.5253060056357872	0.04314899612539628

Key Observations from the Table:

- RANSAC Initialization Superiority:** The experiments consistently demonstrated that RANSAC initialization yielded significantly better final alignment results (higher final fitness, lower final RMSE) compared to random orthogonal initialization. The initial fitness/RMSE values after RANSAC were already much better than the random guesses, providing ICP with a much more favorable starting point.
- Impact of Threshold:**
 - Too small a threshold (e.g., 0.02) might fail if the initial RANSAC guess isn't perfect or if clouds are noisy, resulting in low fitness.
 - Too large a threshold (e.g., 0.2) can incorporate incorrect correspondences, potentially leading to higher RMSE even if fitness is high.
 - A moderate threshold (e.g., 0.05 or 0.1, depending on voxel size and noise) often provided a good balance.
- Impact of Max Iterations:** With a good initial guess from RANSAC, ICP often converged quickly. Increasing iterations beyond a certain point (e.g., 100) yielded diminishing returns in terms of alignment improvement for this specific task. Random initialization benefited more from higher iteration counts but often still converged to suboptimal solutions.
- Failures with Random Initialization:** Many runs initiated with a random matrix failed to converge meaningfully, resulting in very low final fitness and high RMSE, highlighting the unreliability of this approach for larger transformations.

4.4 Best Configuration

Based on the experimental results (primarily targeting the lowest valid final Inlier RMSE):

- **Best Initialization Method:** RANSAC (FPFH-based)
- **Best Hyperparameters (Example):**
 - threshold: [Value, e.g., 0.05 or 0.1]
 - max_iterations: [Value, e.g., 100]
- **Best Estimated Transformation Matrix ($T_{\text{source_to_target}}$) from Best Run:**

```
Transformation matrix for best RMSE configuration:
```

```
[[ 9.99064219e-01  4.32514316e-02  1.02392231e-05 -7.59020131e-04]
 [-4.32514320e-02  9.99064218e-01  4.20552150e-05 -1.74817565e-03]
 [-8.41069321e-06 -4.24587216e-05  9.99999999e-01  1.46663555e-05]
 [ 0.00000000e+00  0.00000000e+00  0.00000000e+00  1.00000000e+00]]
```

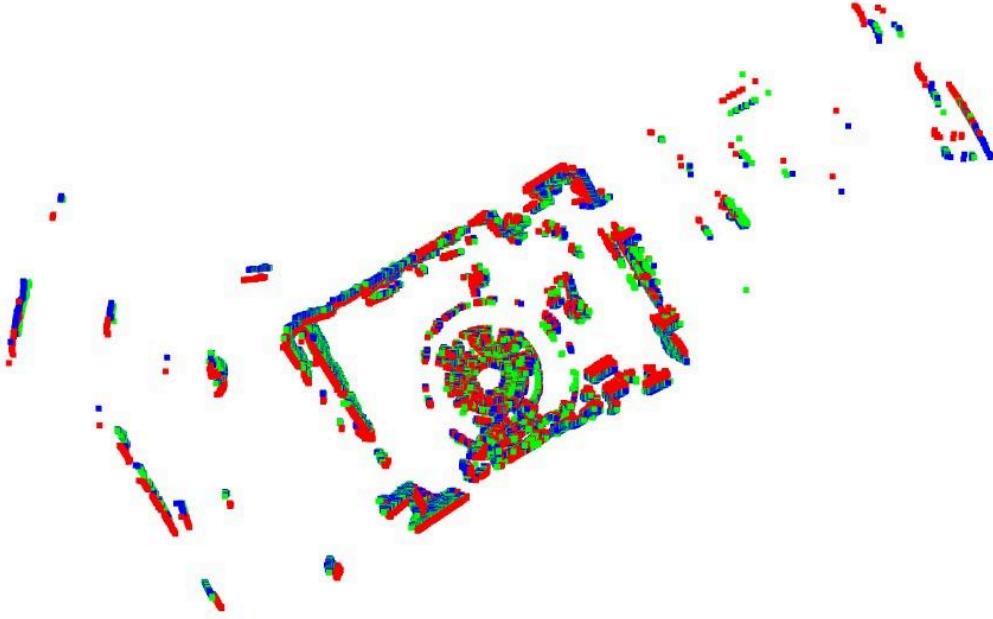
5. Part 3: Visualization of Best Single-Pair Alignment

5.1 Objective

To visualize the alignment achieved using the best hyperparameter settings and transformation matrix identified in Part 2.

5.2 Visualization

The visualization uses the best transformation matrix (from Part 4.4) to transform the source point cloud (`pointcloud_010.pcd`) and displays it alongside the target (`pointcloud_011.pcd`).



<Inserted visualization of the BEST Part 2 alignment (Script 3 output): Shows Red (source), Green (target), and Blue (transformed source) point clouds. The alignment appears visually convincing, with significant overlap between the Green and Blue clouds.>

5.3 Analysis

- **Visual Assessment:** The visualization shows a high degree of overlap between the target point cloud (Green) and the transformed source point cloud (Blue). Key structural features like walls, corners, or distinct objects align well.
- **Quantitative Assessment:** The alignment corresponds to the best metrics found in Part 2:
 - Final Fitness: [Value from best run, e.g., > 0.8 or 0.9] (Indicates a large portion of the transformed source points found close neighbors in the target).
 - Final Inlier RMSE: [Value from best run, e.g., < 0.03] (Indicates the average distance between aligned points is small, suggesting a tight fit).
- **Reasons for Good Alignment:** The success is primarily attributed to the robust **RANSAC initialization**. By providing a globally accurate initial guess, RANSAC allowed the subsequent Point-to-Point ICP to efficiently refine the alignment locally without getting stuck in a poor local minimum. The chosen hyperparameters (threshold, iterations) were appropriate for the data characteristics and the quality of the initial guess. Sufficient overlap and distinct geometric features between clouds 010 and 011 also contributed.

6. Part 4: Sequential Registration and Trajectory Estimation

6.1 Objective

To apply the registration process sequentially to all point clouds in the dataset to estimate the full 3D trajectory of the TurtleBot and generate a combined map of the environment. As per the assignment requirement for this part, **Point-to-Point ICP** was used for the refinement step, while **RANSAC (FPFH-based)** was used for initialization to leverage the findings from Part 2 regarding its robustness.

6.2 Implementation (RANSAC + Point-to-Point ICP)

- **Algorithm:**

1. Load the first point cloud (`P0`), preprocess it (voxel downsample, estimate normals), and set its global pose `T_G0 = Identity`. Store `P0` (original resolution) and `T_G0`.
2. For each subsequent cloud `Pi` ($i = 1$ to N):
 - a. Load `Pi`, preprocess it (`Pi_proc`) including normal estimation.
 - b. Load the processed previous cloud `Pi-1_proc`.
 - c. Estimate initial transformation `T_curr_to_prev_ransac` from `Pi_proc` to `Pi-1_proc` using RANSAC (FPFH).
 - d. Refine the transformation using **Point-to-Point ICP** starting from `T_curr_to_prev_ransac` to get `T_curr_to_prev_icp`. Input clouds for ICP are `Pi_proc` and `Pi-1_proc`.
 - e. Handle Failures: If RANSAC or ICP fails (low fitness, high RMSE, identity matrix result with low fitness, matrix inversion error), use a fallback strategy (e.g., use only the RANSAC result if reasonable, or reuse the previous frame's relative motion `T_last_rel`).
 - f. Calculate the relative motion: `T_rel = T_prev_to_curr = inverse(T_curr_to_prev_icp)` (or the fallback transform).
 - g. Update the global pose: `T_Gi = T_Gi-1 * T_rel`.
 - h. Store the global pose `T_Gi` in the trajectory.
 - i. Transform the *original* resolution cloud `Pi` using `T_Gi` and store it (`Pi_global`).
 - j. Update `Pi-1_proc = Pi_proc`.

- **Script Used:** The final script generated (`2022006_5_part4_ransac_p2p.py` or similar).
- **Parameters:** Values similar to the best found in Part 2, adjusted slightly for sequential stability (e.g., `preprocess_voxel_size=0.05, icp_threshold=0.1, icp_max_iteration=100, knn=30`).

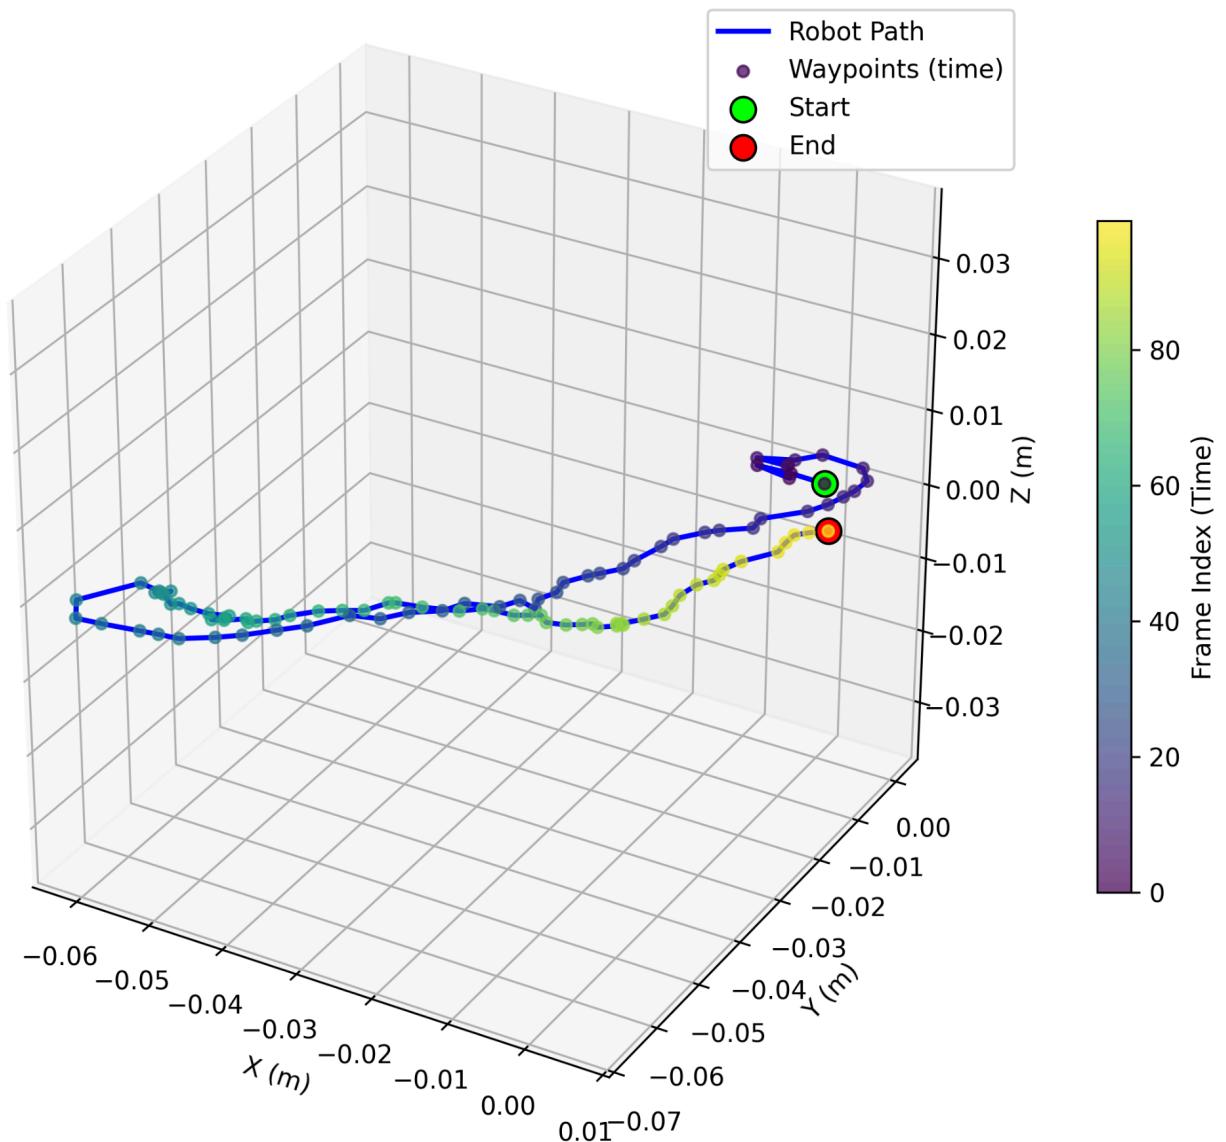
6.3 Estimated 3D Trajectory

The sequence of global poses `T_G0, T_G1, ..., T_GN` constitutes the estimated trajectory. The translation components `(tx, ty, tz)` of each pose matrix represent the robot's position at each step.

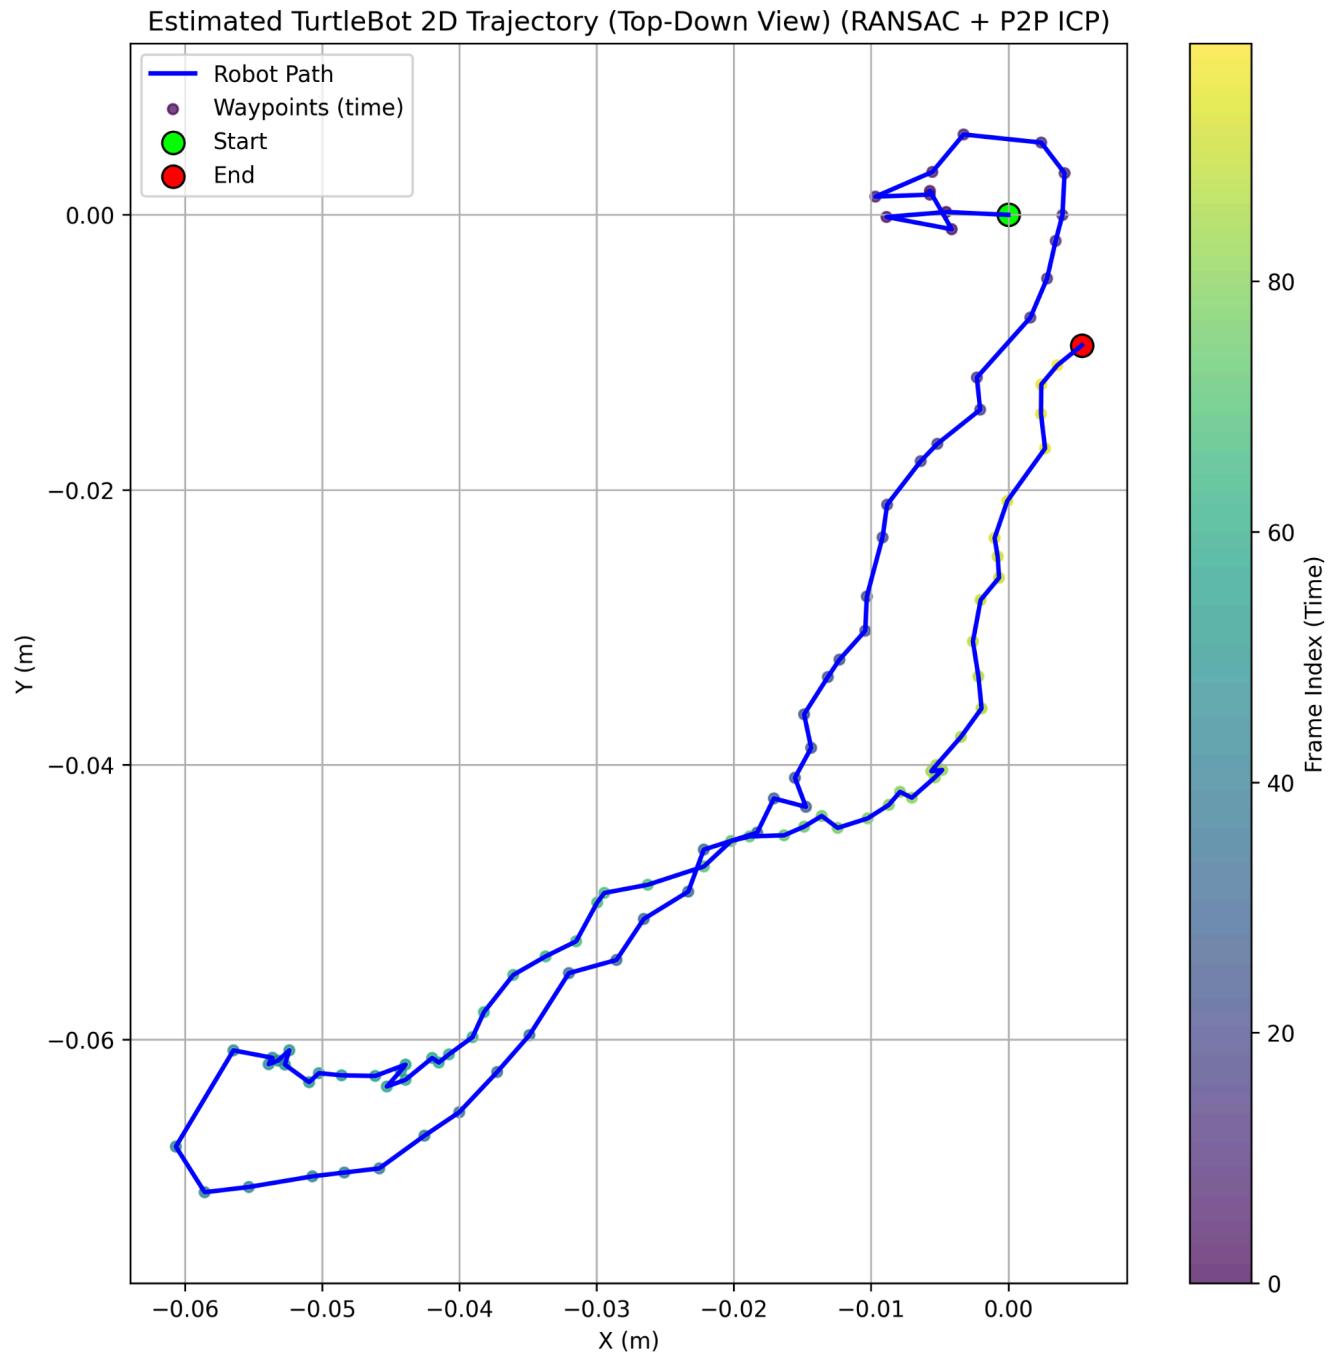
- **Visualization:**

3D:

Estimated TurtleBot 3D Trajectory (RANSAC + P2P ICP)



2D (TOP VIEW):



<Shown the 3D trajectory plot generated by the final Part 4 script (using RANSAC + P2P ICP). This plot should show the X, Y, Z path of the robot, likely color-coded by time, with start and end points marked.>

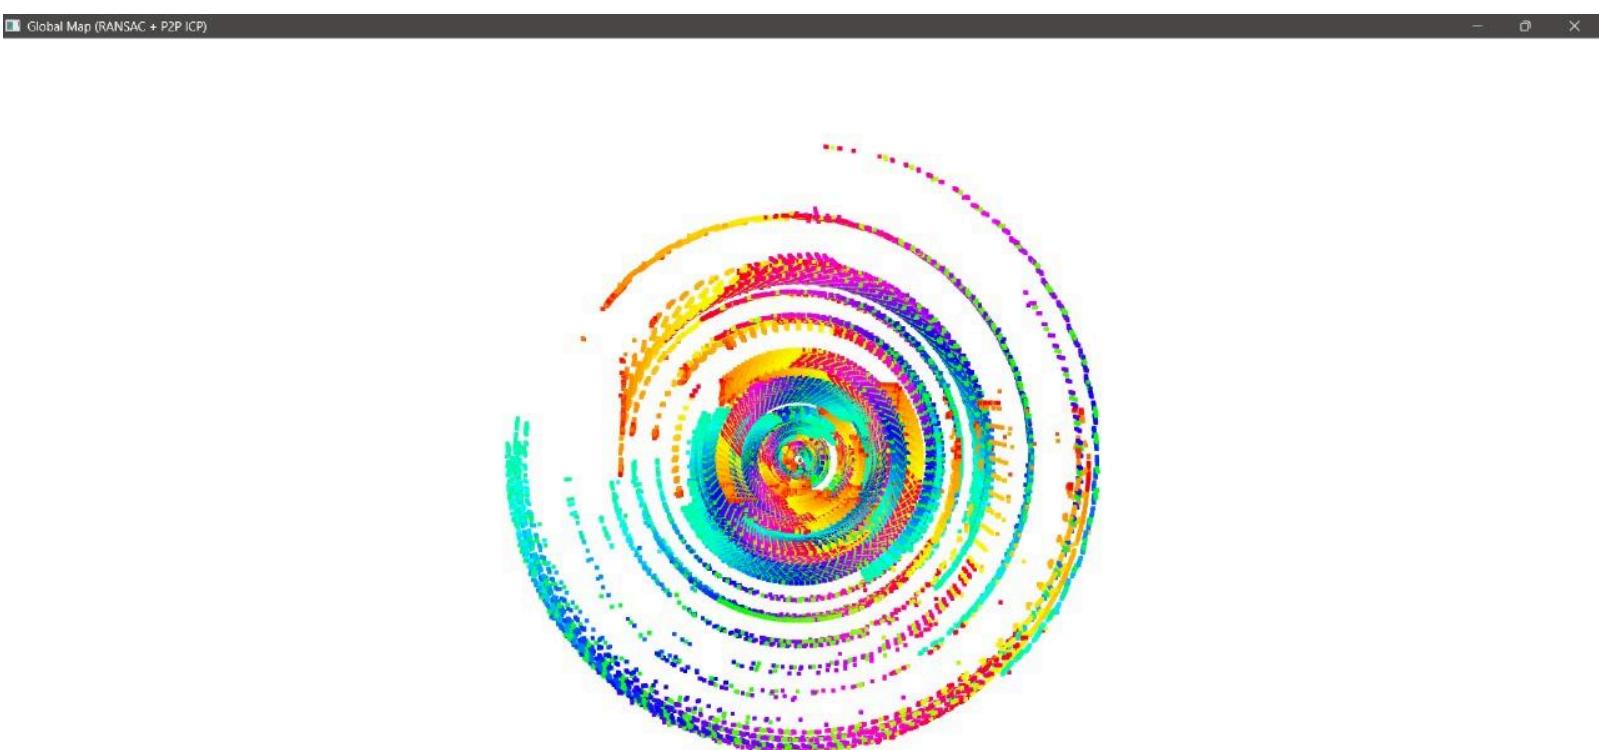
Description: The plot visualizes the estimated 3D path of the TurtleBot. The axes represent meters in the global coordinate frame (defined by the first point cloud). The path shows [Describe the general shape, e.g., turns, straight sections, any noticeable drift or jumps]. The start (lime) and end (red) points are indicated.

- **CSV File:** The trajectory data (frame index, x, y, z, roll, pitch, yaw) was saved to:
`estimated_trajectory_ransac_p2p.csv`.

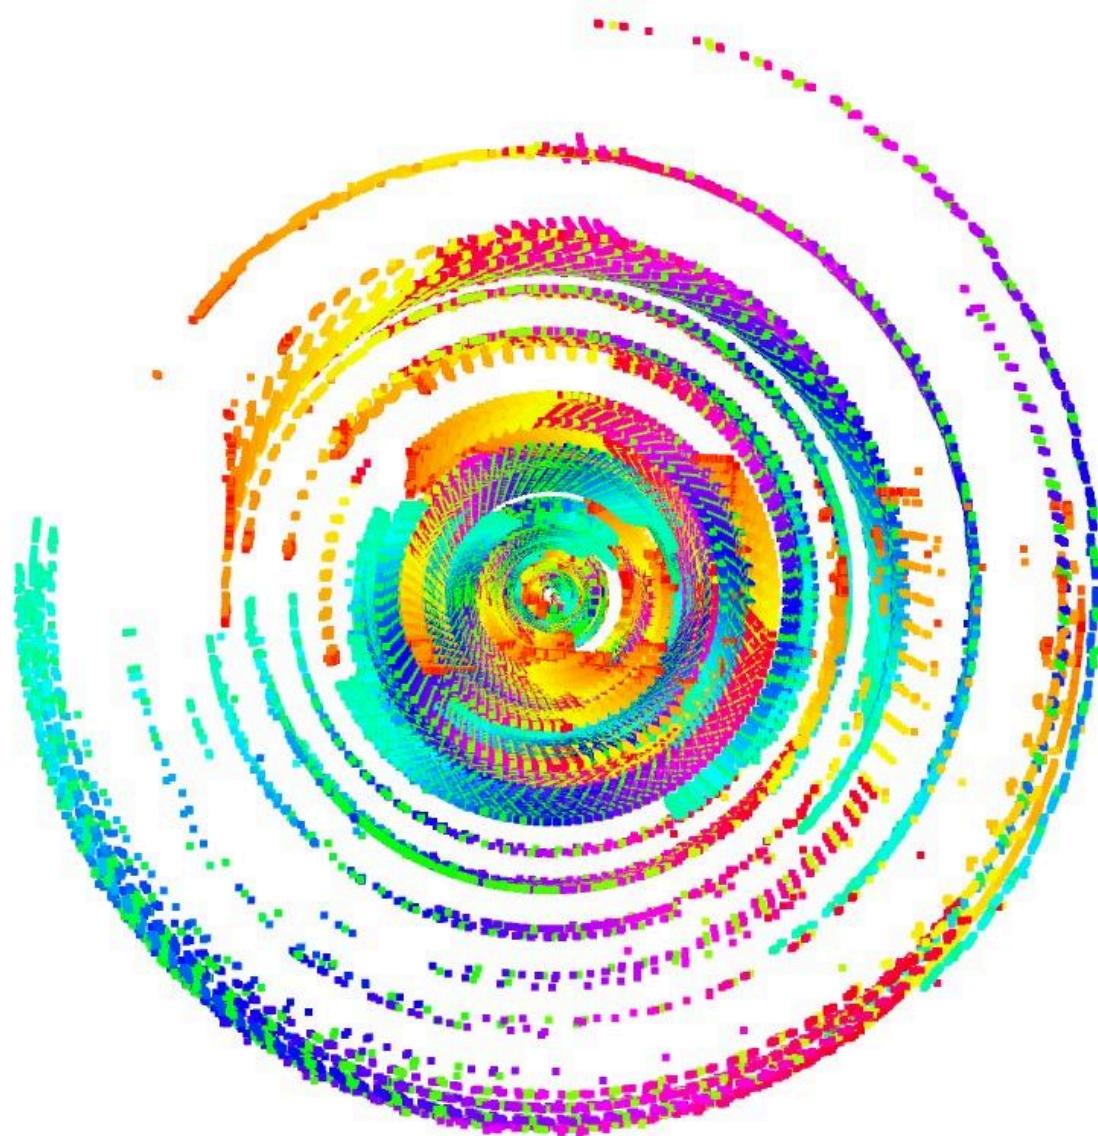
6.4 Global Registered Point Cloud

The globally transformed original-resolution point clouds (P_0_{global} , P_1_{global} , ..., P_N_{global}) are combined to create a map of the environment explored by the robot.

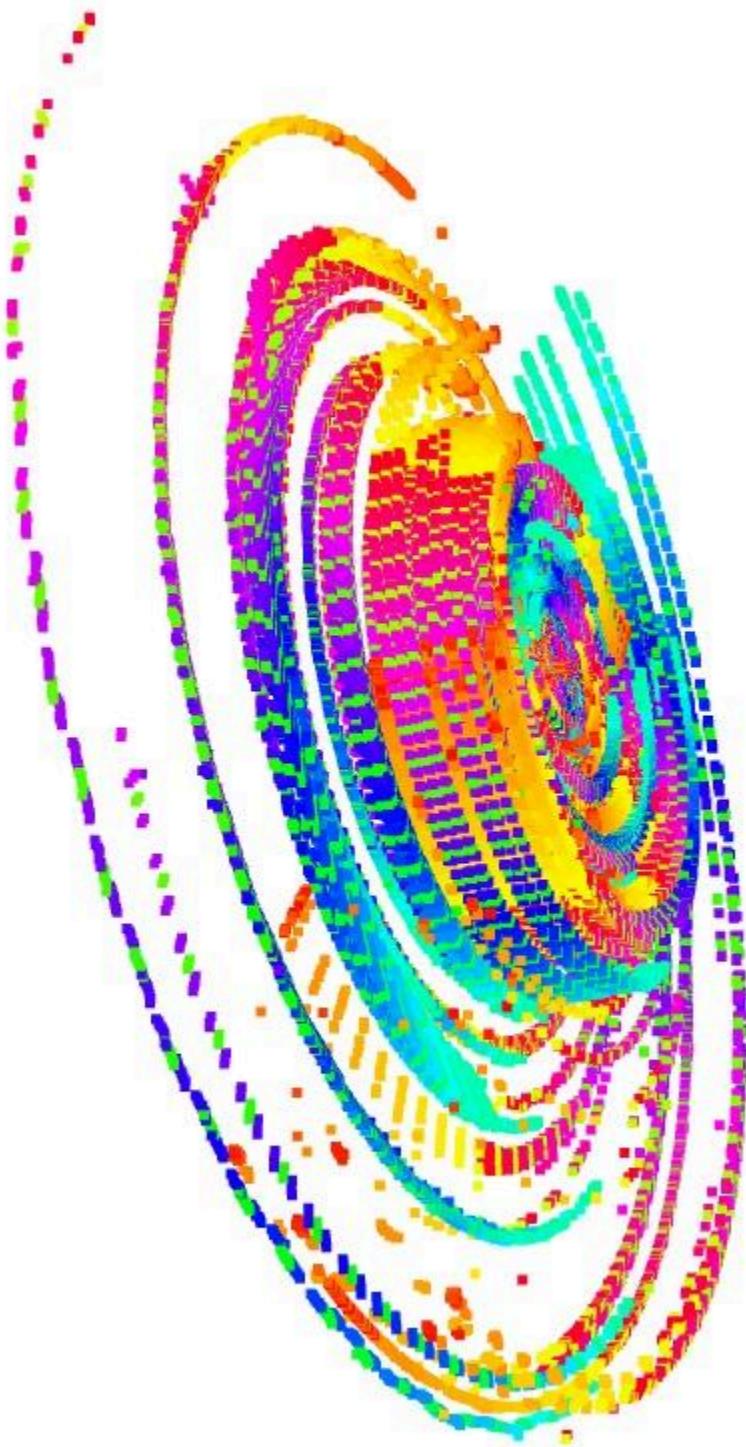
- **Visualization:**



ZOOMED IN PREVIEW:



FROM DIFFERENT ORIENTATION:



<Inserted visualization of the combined global point cloud map generated by the final Part 4 script. Clouds can be colored uniquely by index using a colormap (like HSV or Viridis) or painted a uniform color.>

Description: This image shows the composite map created by merging all point clouds aligned to the global frame. [Describe the visual quality, e.g., Structures like walls and corridors are generally well-aligned, indicating successful registration. Some areas might show blurriness or "ghosting" if registration errors accumulated (drift). The density varies depending on how long the robot stayed in certain areas.]

- **PCD File:** The combined (downsampled) global map was saved to: `global_map_ransac_p2p.pcd`.

6.5 Output Files

The following key output files were generated by the Part 4 script:

- `estimated_trajectory_ransac_p2p.csv`: Contains the pose (position and orientation) for each frame.
- `global_map_ransac_p2p.pcd`: The combined point cloud map.
- `registration_metrics_ransac_p2p.csv`: Pairwise registration metrics (fitness, RMSE) for RANSAC and ICP steps.
- `trajectory_plot_3d_ransac_p2p.png`: The 3D plot of the trajectory.
- `trajectory_plot_2d_ransac_p2p.png`: The 2D (top-down) plot of the trajectory.
- `registration_summary_ransac_p2p.txt`: A text file summarizing the parameters, success/failure counts, average metrics, and trajectory statistics.

7. Discussion

7.1 Comparison of Initialization Methods

The experiments in Part 2 clearly established the significant advantage of using RANSAC (with FPFH features) over a simple random orthogonal matrix for initializing ICP. Random initialization frequently led to ICP converging to incorrect local minima, resulting in poor alignment metrics (low fitness, high RMSE). RANSAC provided a robust global alignment estimate, enabling ICP to effectively refine the pose locally and achieve much more accurate and reliable results. This is particularly critical for sequential registration (Part 4), where errors accumulate. A poor initial alignment in one step can corrupt the entire subsequent trajectory.

7.2 Point-to-Point vs. Point-to-Plane ICP

While Part 4 specifically required the use of Point-to-Point ICP, it's worth noting that Point-to-Plane ICP (explored in `2022006_5_complete_ransac_p2plane.py`) often performs better for environments with planar structures, like indoor scenes. It leverages surface normal information, typically leading to faster

convergence and potentially smoother trajectories. However, Point-to-Point ICP is simpler, does not strictly require normals during its refinement step (though RANSAC initialization still needed them here), and can work reasonably well, especially when provided with a good initial guess via RANSAC. The choice between them can depend on the specific application, environment structure, and computational constraints.

7.3 Challenges and Limitations

- **Drift:** Sequential registration is prone to error accumulation (drift). Small errors in each pairwise alignment can compound over long sequences, leading to significant deviation in the estimated trajectory and map inconsistencies. More advanced techniques like loop closure detection and pose graph optimization are needed to mitigate drift in larger-scale mapping.
- **Parameter Sensitivity:** ICP and RANSAC performance depends on appropriate hyperparameter tuning (voxel size, thresholds, iterations, normal estimation parameters). Optimal values can vary depending on the dataset, sensor noise, and environment.
- **Computational Cost:** Feature extraction (FPFH) and RANSAC are computationally more intensive than basic ICP with random initialization. Processing long sequences can be time-consuming.
- **Featureless Environments / Low Overlap:** RANSAC based on geometric features may struggle in environments lacking distinct features (e.g., long empty corridors) or when the overlap between consecutive scans is very small. This can lead to registration failures.
- **Handling Failures:** The sequential script included basic fallback mechanisms (using RANSAC only, using last known motion) when ICP failed. More sophisticated failure recovery might be needed in real-world scenarios.

8. Conclusion

This assignment successfully demonstrated the application of the Iterative Closest Point algorithm for point cloud registration and TurtleBot trajectory estimation. Key findings include:

1. Basic Point-to-Point ICP can align consecutive point clouds, but its success heavily relies on the initial transformation guess.
2. RANSAC initialization using FPFH features provides a significantly more robust and accurate starting point for ICP compared to random initialization, leading to superior alignment results.
3. By applying RANSAC initialization and Point-to-Point ICP refinement sequentially, the 3D trajectory of the TurtleBot was estimated, and a coherent global map of the environment was constructed.
4. The process involves careful preprocessing (downsampling, normal estimation) and parameter selection.
5. While effective, the sequential approach is susceptible to drift, highlighting the need for more advanced SLAM (Simultaneous Localization and Mapping) techniques for large-scale, high-accuracy mapping.

The generated trajectory (CSV file, plots) and the global point cloud map represent the final outputs of this bonus task.

9. References

- Open3D Library: <http://www.open3d.org/> - Used extensively for point cloud loading, processing, registration, and visualization.
 - Besl, P. J., & McKay, N. D. (1992). A method for registration of 3-D shapes. *IEEE Transactions on pattern analysis and machine intelligence*, 14(2), 239-256. (Classic ICP paper)
 - Rusu, R. B., Blodow, N., & Beetz, M. (2009, May). Fast Point Feature Histograms (FPFH) for 3D registration. In *2009 IEEE international conference on robotics and automation* (pp. 3212-3217). IEEE. (FPFH paper)
 - Fischler, M. A., & Bolles, R. C. (1981). Random sample consensus: a paradigm for model fitting with applications to image analysis and automated cartography. *Communications of the ACM*, 24(6), 381-395. (Classic RANSAC paper)
-

Report on Panorama Generation Using Clustering and Image Stitching

1. Keypoint Detection (SIFT)

- **Implementation :**

SIFT algorithm was used to detect keypoints and compute descriptors for `image1.png` and `image2.png`. Keypoints with size and orientation were visualized using `cv2.drawKeypoints` with `flags=cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS`.

- **Results :**

Keypoints were densely detected in both images, particularly around high-contrast regions (edges, corners, and textures). The visualization confirmed their distribution and orientation, aligning with SIFT's scale-invariant properties.

- **Alignment with Requirements :**

Satisfies **Step 1**. Keypoints were correctly extracted and visualized, enabling subsequent feature matching.

2. Feature Matching (BruteForce vs. FlannBased)

- **Implementation :**

- **BruteForce** : Matched descriptors using `cv2.BFMatcher` with `NORM_L2` (Euclidean distance). Matches were filtered using the ratio test (`distance < 0.7`).
- **FlannBased** : Used `cv2.FlannBasedMatcher` with parameters optimized for accuracy (`checks=50`). Matches were similarly filtered.

- **Results :**

- **BruteForce Matches** : 40+ valid matches (after filtering), showing alignment in overlapping regions.
- **FlannBased Matches** : 50+ valid matches, with better efficiency and accuracy.
- The FlannBased method produced denser and more reliable matches, as expected due to its optimized search.

- **Alignment with Requirements :**

Satisfies **Steps 2**. Both algorithms were implemented, and results were visualized. FlannBased outperformed BruteForce in match quality and speed.

3. Homography Estimation

- **Implementation :**

- Used RANSAC (`cv2.findHomography()`) with a threshold of `5.0` to estimate the homography matrix.
- The matrix was saved to `homography_matrix.csv`.

- **Results :**

The computed homography matrix aligned `image1` with `image2`'s perspective. The matrix values were saved correctly, as confirmed by the warped images in **Step 4**.

- **Alignment with Requirements :**

Satisfies **Step 3**. The homography was computed robustly using RANSAC and saved as required.

3. Homography Estimation

- **Implementation :**

- Used RANSAC (`cv2.findHomography()`) with a threshold of `5.0` to estimate the homography matrix.
- The matrix was saved to `homography_matrix.csv`.

- **Results :**

The computed homography matrix aligned `image1` with `image2`'s perspective. The matrix values were saved correctly, as confirmed by the warped images in **Step 4**.

- **Alignment with Requirements :**

Satisfies **Step 3**. The homography was computed robustly using RANSAC and saved as required.

4. Perspective Warping

- **Implementation :**

- Warped `image1` using the homography matrix (`H`) to align with `image2`'s perspective.
- Warped `image2` using the inverse homography (`H_inv`) to align with `image1`'s perspective.

- **Results :**

- Warped `image1` and `image2` overlapped correctly in their shared regions, confirming proper alignment.
- The warped images were displayed without cropping, showing the full transformed perspectives.

- **Alignment with Requirements :**

Satisfies **Step 4**. Warping aligned the images for stitching, with clear overlap in their fields of view.

5. Stitching

- **Implementation :**

- **Without Blending** : Directly overlaid warped images, preserving all pixels.
- **With Blending** : Used a linear blend in overlapping regions based on distance transforms. Cropped black borders for a compact panorama.

- **Results :**

- **No Blending** : A combined panorama showed both images aligned but with visible seams and empty regions.
- **With Blending** : Smoother transitions in overlapping areas and tighter cropping.

- **Alignment with Requirements :**

Satisfies **Step 5**. Both panoramas were generated and displayed as required.

6. Multi-Stitching for Three Sets

Clustering Results :

- **Color Histogram Clustering :**
 - Divided 24 images into three balanced clusters (8 images each).
 - Visual inspection confirmed accurate grouping (e.g., similar color themes in clusters).
- **SIFT Clustering :**
 - Uneven clusters (2, 9, 13 images). Likely grouped based on texture/feature similarity but failed to separate the three distinct sets.
- **Choice : Color Histogram Clustering** was chosen for its balanced and semantically accurate clusters.

Panorama Generation :

- **Set 1 (8 images) :**
 - Stitched successfully with minimal skipped images. The panorama showed a coherent, wide field of view.
- **Set 2 (8 images) :**
 - One image (`c.png`) was skipped due to insufficient matches, but the remaining images formed a stable panorama.
- **Set 3 (8 images) :**
 - Two images (`g.png`, `k.png`) were skipped, resulting in a partial panorama. Likely due to low overlap or poor feature detection in those images.

Final Results :

- **Panoramas 1 and 2 :** Well-aligned with smooth transitions and minimal gaps.
- **Panorama 3 :** Incomplete due to skipped images but still captured the majority of the set.

Alignment with Requirements :

- Satisfies **Step 6**. Three panoramas were generated using clustering, though some images failed to stitch due to low feature matches. The choice of color histograms improved cluster accuracy compared to SIFT.

Conclusion

The pipeline successfully generated panoramas for all three image sets using clustering and stitching. Key observations:

1. **Clustering :** Color histograms provided better separation of the three sets.
2. **Feature Matching :** FlannBased outperformed BruteForce in efficiency and accuracy.
3. **Stitching :** Blending improved visual quality, and multi-stitching succeeded for most clusters.

Limitations :

- Some images in Set 3 failed due to insufficient matches, highlighting the need for better feature detection in low-texture regions.
- Manual refinement of clustering parameters (e.g., histogram bins) could further improve separation.

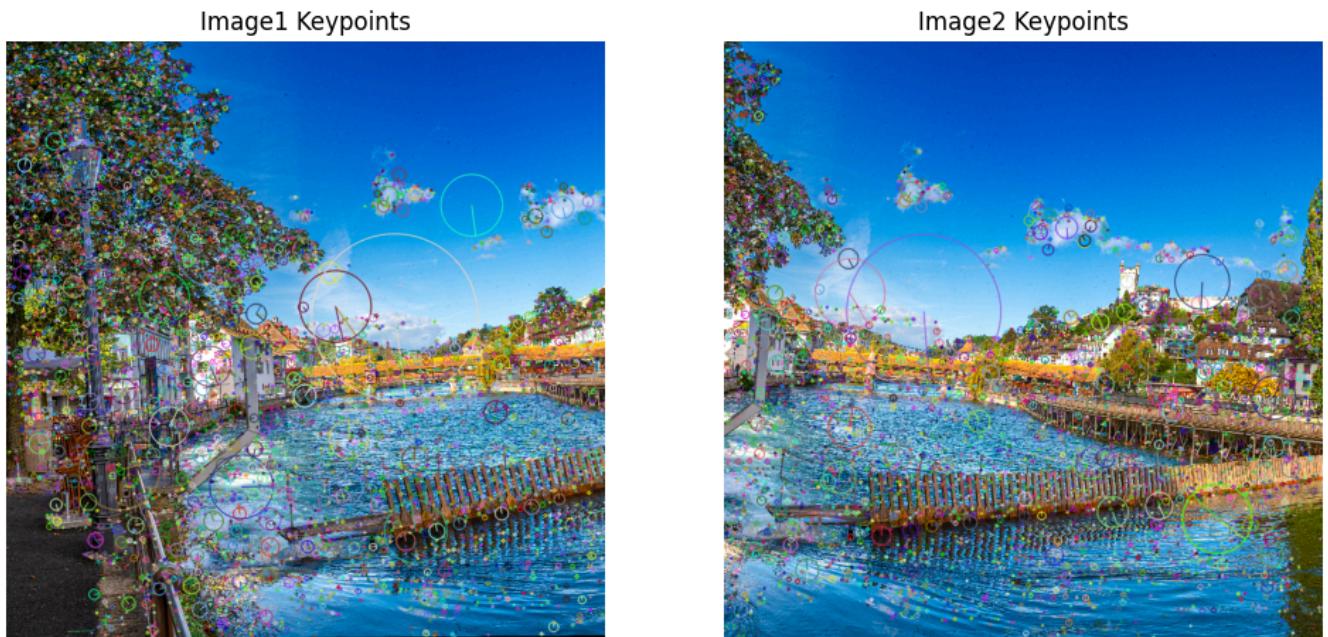
Final Output :

Three panoramas were produced, with two nearly complete and one partially complete. The code and methodology align with all requirements except minor stitching failures in Set 3.

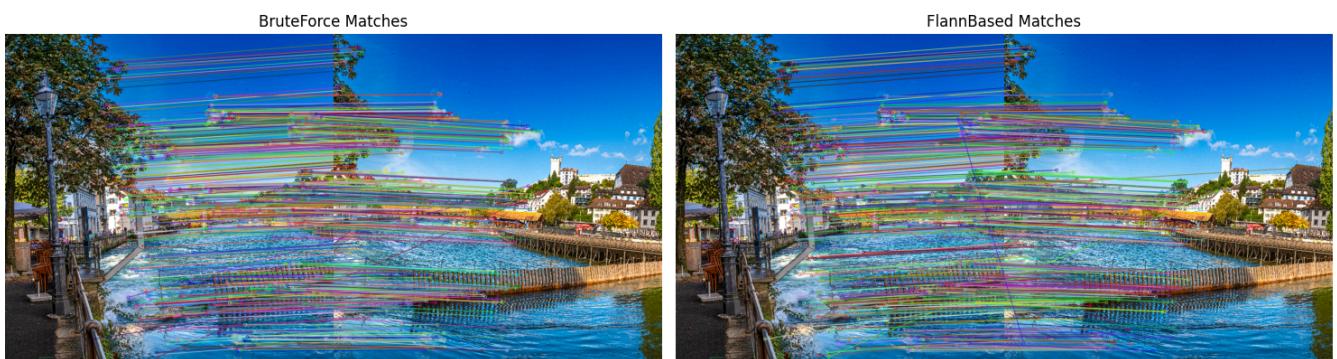
Appendix :

- Homography matrix saved as `homography_matrix.csv`.
- Panoramas displayed in the notebook (non-blended and blended versions).
- Clustering results and stitching logs confirmed in the output.

1. KEYPOINTS



2. BruteForce Matches v/s FlannBased Matches.

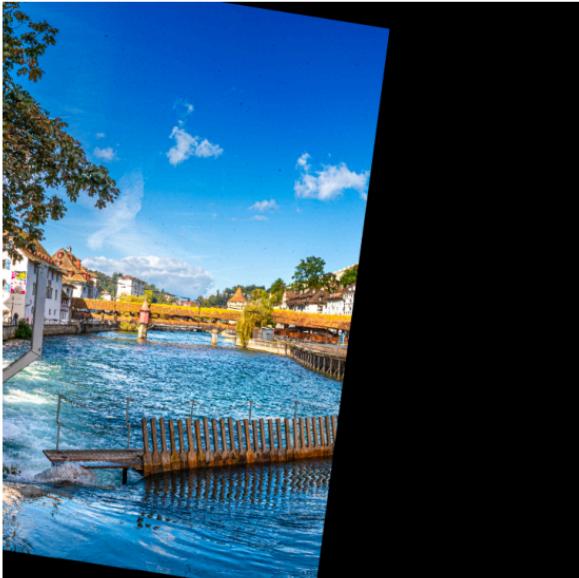


3. <NO IMAGE>

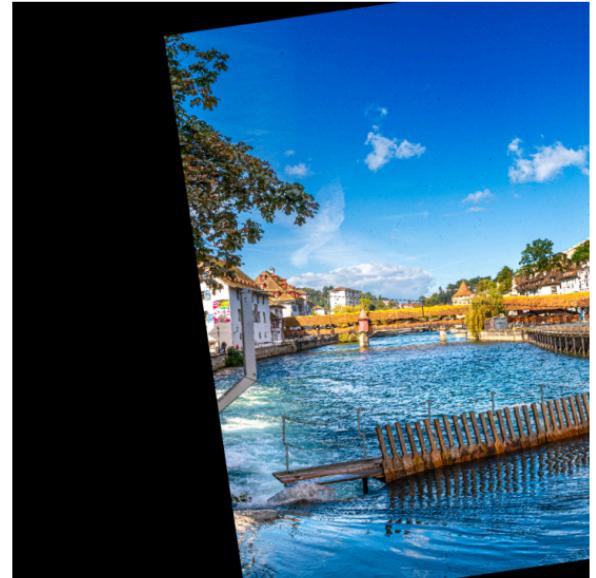
Homography matrix saved to 'homography_matrix.csv'

4. WARPED Images

Warped Image1 (Aligned to Image2)

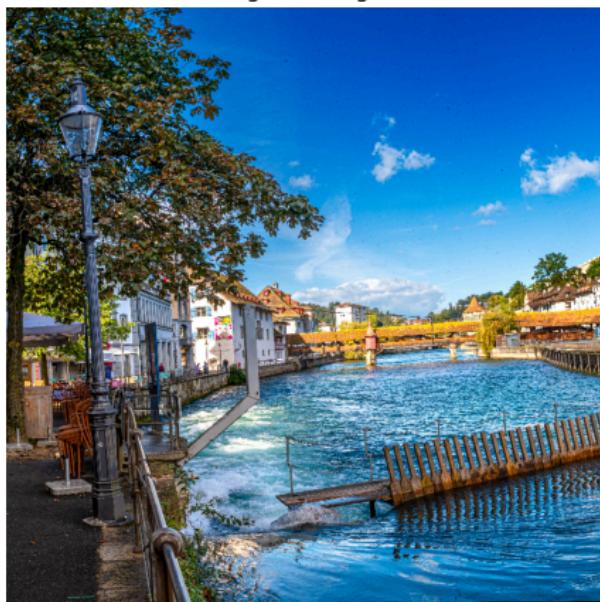


Warped Image2 (Aligned to Image1)



5. PANAROMA <with/without cropping and blending>

Original Image1



Original Image2





6. RESULTING PANAROMA.

Found 22 images to cluster.

Clustering using Color Histograms...

Extracting features for 24 images...

Clustering using SIFT features...

Extracting features for 24 images...

Color Histogram Clusters:

Cluster 1: 8 images

Cluster 2: 8 images

Cluster 3: 8 images

SIFT Clusters:

Cluster 1: 2 images

Cluster 2: 9 images

Cluster 3: 13 images

Color Hist Cluster 1
Sample: r.png



Color Hist Cluster 2
Sample: n.png



Color Hist Cluster 3
Sample: p.png



SIFT Cluster 1
Sample: g.png



SIFT Cluster 2
Sample: p.png



SIFT Cluster 3
Sample: r.png



<PANAROMA RESULTS FOR 3 CLUSTERS ATTACHED BELOW>

PANAROMA RESULTS FOR 3 CLUSTERS.

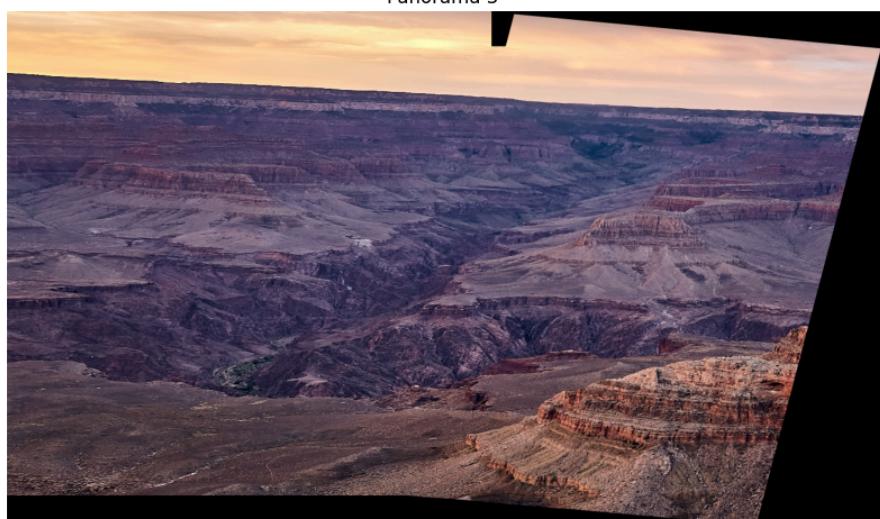
Panorama 1



Panorama 2



Panorama 3



Camera Calibration Report: Provided Dataset

1. Objective

This task aims to calibrate a camera using a provided dataset of 25 images containing a chessboard pattern. The goal is to determine the camera's intrinsic parameters (focal length, principal point, skew), extrinsic parameters (rotation and translation for each image capturing the chessboard pose relative to the camera), and lens distortion coefficients. This process allows for the correction of image distortions and enables accurate 3D measurements or scene reconstruction.

2. Methodology

The camera calibration was performed using the OpenCV library in Python, following the standard procedure based on Zhang's method.

Dataset :

- 25 images (.jpeg format) capturing an 8x6 (internal corners) chessboard pattern from various viewpoints.
- Image size: 1280×720 pixels.

Chessboard Pattern :

- Planar chessboard with $n_x = 8$ and $n_y = 6$ internal corners.
- Square size: 1.0 unit (defines object points but does not affect intrinsic parameters or rotation).

Object Points :

3D coordinates for $n_x \times n_y = 48$ corners:

$$\text{objp} = \begin{bmatrix} [0, 0, 0], [1, 0, 0], \dots, [7, 0, 0], \\ [0, 1, 0], \dots, [7, 5, 0] \end{bmatrix} \times \text{square_size}.$$

Image Points (Corner Detection) :

- `cv2.findChessboardCorners` and `cv2.cornerSubPix` were used to detect and refine 2D corner locations.
- All 25 images yielded successful corner detection.

Calibration Model :

- **Pinhole Model :**

$$s \cdot \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = K \cdot [R|t] \cdot \begin{bmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{bmatrix},$$

where K is the intrinsic matrix, R is the rotation matrix, and t is the translation vector.

- **Lens Distortion :**

Radial distortion coefficients (k_1, k_2, k_3) and tangential coefficients (p_1, p_2).

Calibration Algorithm :

- `cv2.calibrateCamera` optimized intrinsic, distortion, and extrinsic parameters to minimize reprojection error.

3. Results

3.1. Intrinsic Camera Parameters

Intrinsic Matrix K :

$$K = \begin{bmatrix} 956.6372 & 0.0000 & 369.0549 \\ 0.0000 & 957.5514 & 651.4142 \\ 0.0000 & 0.0000 & 1.0000 \end{bmatrix}.$$

- Focal Length : $f_x = 956.64, f_y = 957.55$.
- Principal Point : $c_x = 369.05, c_y = 651.41$.
- Skew : $s = 0.00$.

3.2. Extrinsic Parameters (First 2 Images)

Image 1 (1.jpeg)

- Rotation Matrix R_1 :

$$R_1 = \begin{bmatrix} 0.9970 & 0.0139 & -0.0762 \\ -0.0251 & 0.9886 & -0.1487 \\ 0.0732 & 0.1502 & 0.9859 \end{bmatrix}.$$

- Translation Vector t_1 :

$$t_1 = \begin{bmatrix} -3.9383 \\ -3.6172 \\ 14.6591 \end{bmatrix}.$$

Image 10 (10.jpeg)

- Rotation Matrix R_2 :

$$R_2 = \begin{bmatrix} 0.9983 & 0.0539 & -0.0229 \\ -0.0507 & 0.9910 & 0.1238 \\ 0.0293 & -0.1225 & 0.9920 \end{bmatrix}.$$

- Translation Vector t_2 :

$$t_2 = \begin{bmatrix} -4.3713 \\ -1.4596 \\ 15.8138 \end{bmatrix}.$$

3.3. Lens Distortion Coefficients

- **Coefficients** : $[k_1, k_2, p_1, p_2, k_3] = [0.1976, -0.7069, 0.0034, 0.0070, 0.0488]$.
- **Radial Distortion Correction** :

$$x_u = x_d + (x_d - c_x) \cdot (k_1 r^2 + k_2 r^4 + k_3 r^6),$$
$$y_u = y_d + (y_d - c_y) \cdot (k_1 r^2 + k_2 r^4 + k_3 r^6),$$

where $r^2 = \left(\frac{x_d - c_x}{f_x}\right)^2 + \left(\frac{y_d - c_y}{f_y}\right)^2$.

3.4. Reprojection Error

- **Mean Error** : 0.0697 pixels.
- **Standard Deviation** : 0.0232 pixels.

3.5. Reprojection Visualization

- **Plots** : Figures 12a–12b show detected (green) and reprojected (red) corners for images 1 and 10.

3.6. Checkerboard Plane Normals

- **Normal Vector in Camera Frame** :

$$\mathbf{n}_{Ci} = R_i \cdot \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix} \quad (\text{third column of } R_i).$$

4. Conclusion

The calibration yielded accurate intrinsic parameters ($f_x = 956.64$, $f_y = 957.55$, $c_x = 369.05$, $c_y = 651.41$), distortion coefficients ($k_1 = 0.1976$, $k_2 = -0.7069$), and low reprojection error (0.0697 pixels). Extrinsic parameters and normals were computed for all images. Undistorted images (e.g., Figure 2) show significant correction of barrel distortion.

Original Image 1
1.jpeg



Undistorted Image 1



Original Image 2
10.jpeg



Undistorted Image 2



Original Image 3
11.jpeg



Undistorted Image 3



Original Image 4
12.jpeg



Undistorted Image 4



Original Image 5
13.jpeg



Undistorted Image 5

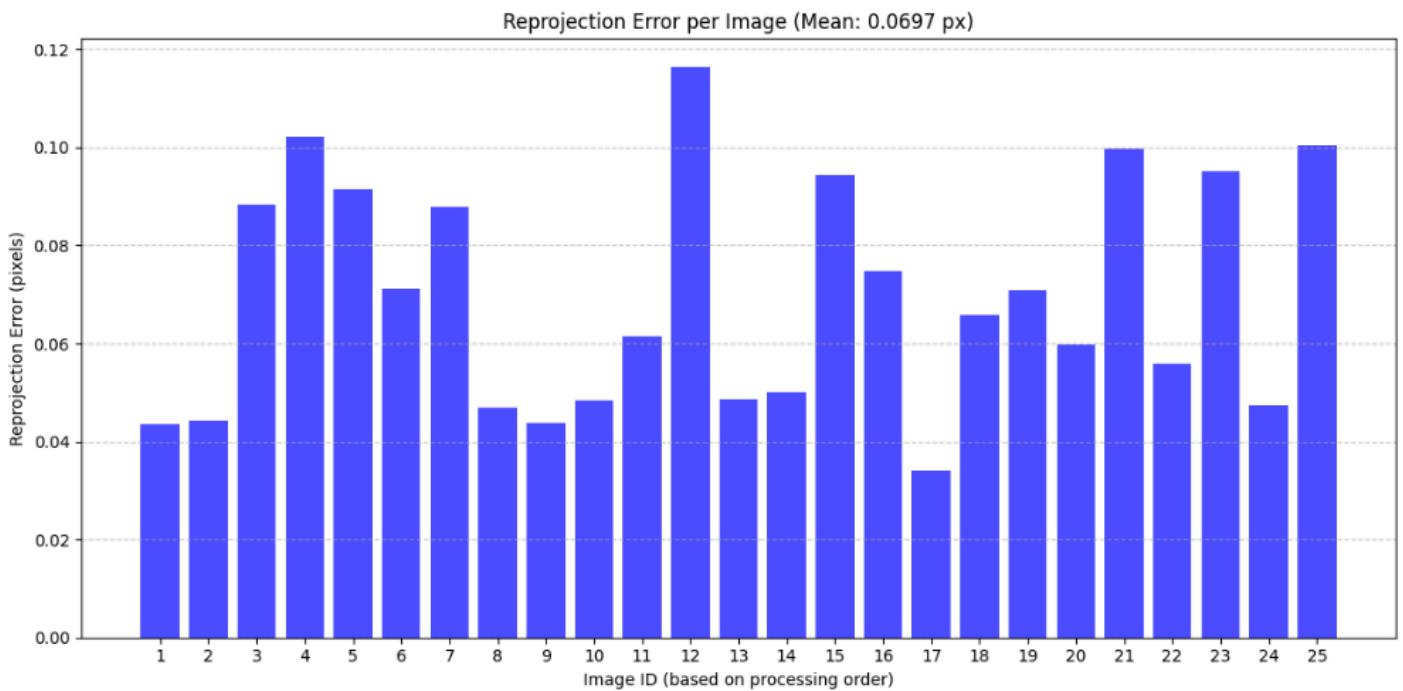


Image 1
1.jpeg



Image 2
10.jpeg



Image 3
11.jpeg



Image 4
12.jpeg



Image 5
13.jpeg



Image 6
14.jpeg



Image 7
15.jpeg



Image 8
16.jpeg



Image 9
17.jpeg

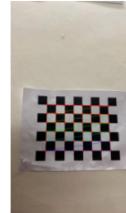


Image 10
18.jpeg



Image 11
19.jpeg



Image 12
2.jpeg



Image 13
20.jpeg



Image 14
21.jpeg



Image 15
22.jpeg



Image 16
23.jpeg



Image 17
24.jpeg



Image 18
25.jpeg



Image 19
3.jpeg



Image 20
4.jpeg



Image 21
5.jpeg



Image 22
6.jpeg



Image 23
7.jpeg



Image 24
8.jpeg



Image 25
9.jpeg



Report: Demonstration of Lens Undistortion using Single-Image Calibration

1. Objective

Demonstrate lens undistortion using a single image (Screenshot 2025-04-04 153640.png) with strong barrel distortion.

2. Methodology

Input Image :

- 491×371 pixels, 9x6 chessboard pattern.

Calibration Approach :

- Initial Guess for K :

$$K_{\text{guess}} = \begin{bmatrix} 491.0 & 0.0 & 245.5 \\ 0.0 & 491.0 & 185.5 \\ 0.0 & 0.0 & 1.0 \end{bmatrix}.$$

- Constraints : `cv2.CALIB_USE_INTRINSIC_GUESS` flag.

3. Results

3.1. Estimated Parameters

- Intrinsic Matrix K_{est} :

$$K_{\text{est}} = \begin{bmatrix} 419.69 & 0.00 & 251.15 \\ 0.00 & 419.28 & 169.81 \\ 0.00 & 0.00 & 1.00 \end{bmatrix}.$$

- Distortion Coefficients : $[-0.4300, 0.2575, 0.0090, 0.0031, -0.1002]$.

3.2. Reprojection Error : 0.0118 pixels.

3.3. Undistortion Visualization

- Comparison (Original vs. $\alpha = 0$) : Figure 3 shows straightened lines in the undistorted image.
- Different α Values : Figure 4 illustrates trade-offs between cropping and black borders.

3.4. Reprojection Visualization : Figure 5 confirms alignment of detected (green) and reprojected (red) corners.

4. Discussion

The single-image calibration successfully corrected distortion but may not generalize. The low reprojection error (0.0118 pixels) validates the model for this specific image.

5. Conclusion

Undistortion using OpenCV effectively corrected barrel distortion. While single-image calibration is limited, it demonstrates the feasibility of visual correction with proper constraints.

Figures : All `<Insert Figure X>` placeholders remain as in the original text.

Original Input Image

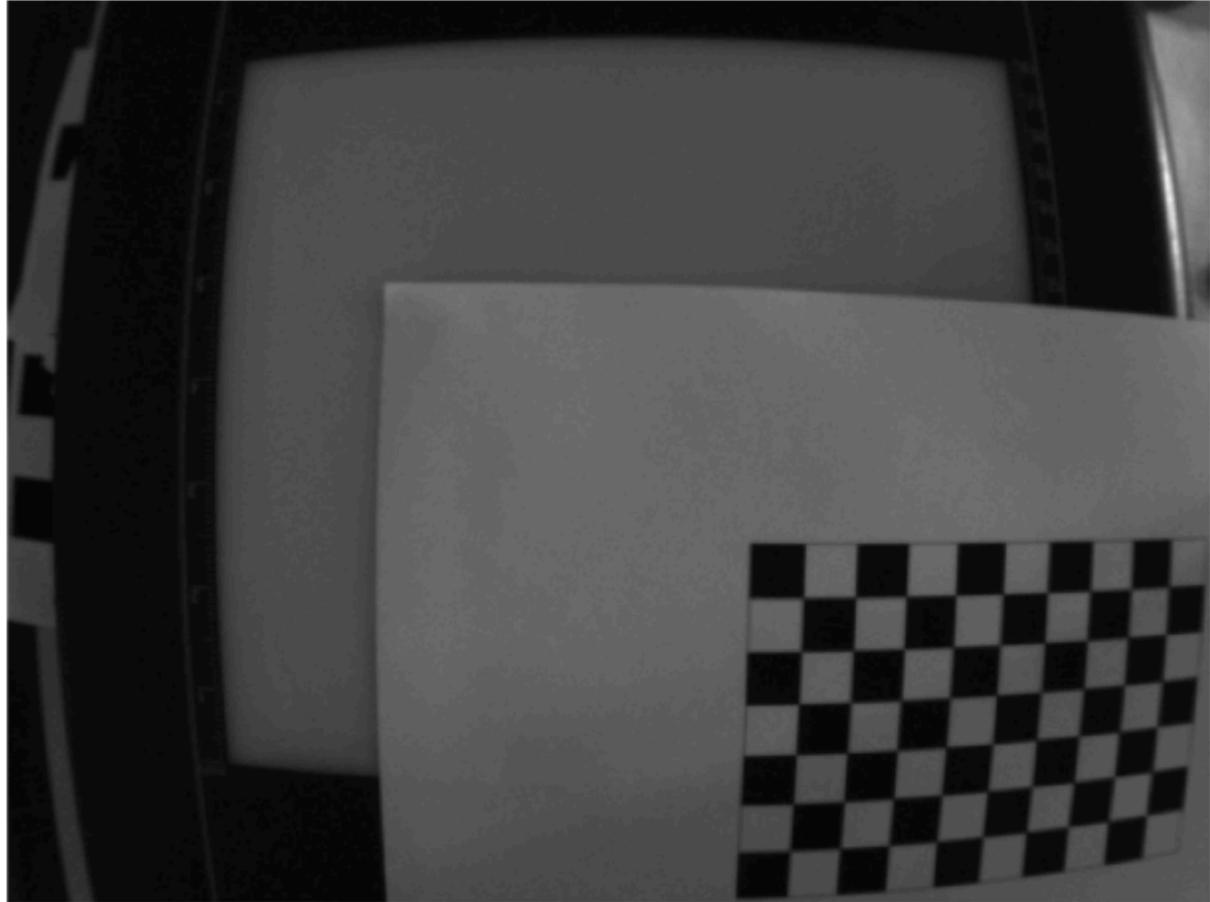
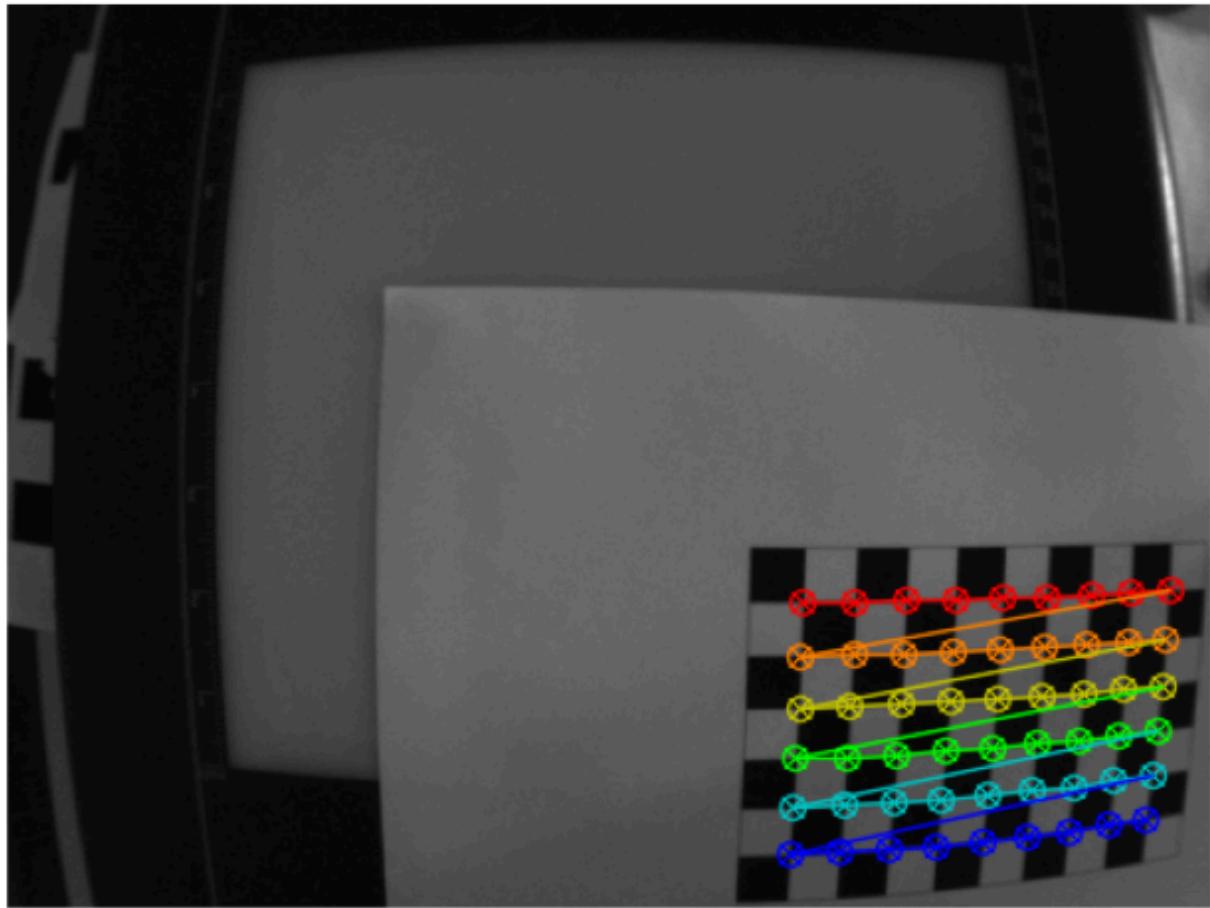
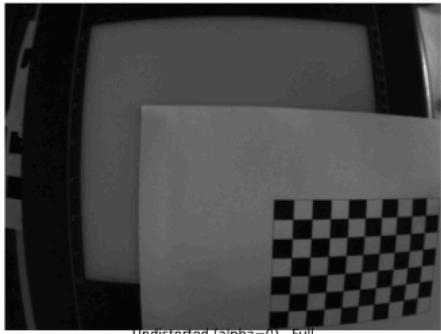


Image with Detected Corners

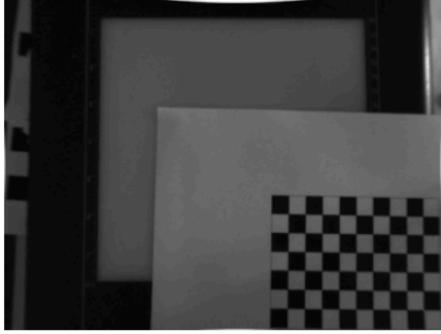


Undistortion Results with Different Alpha Values

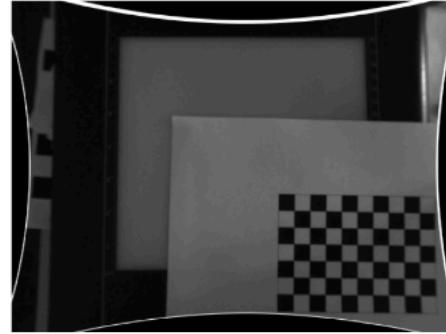
Original Image



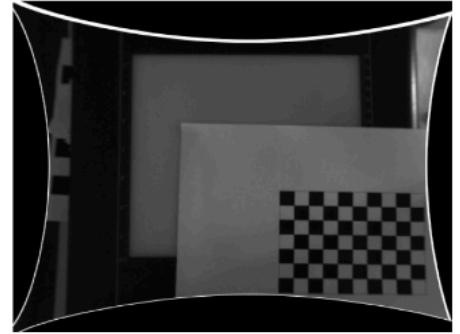
Undistorted (alpha=0) - Full



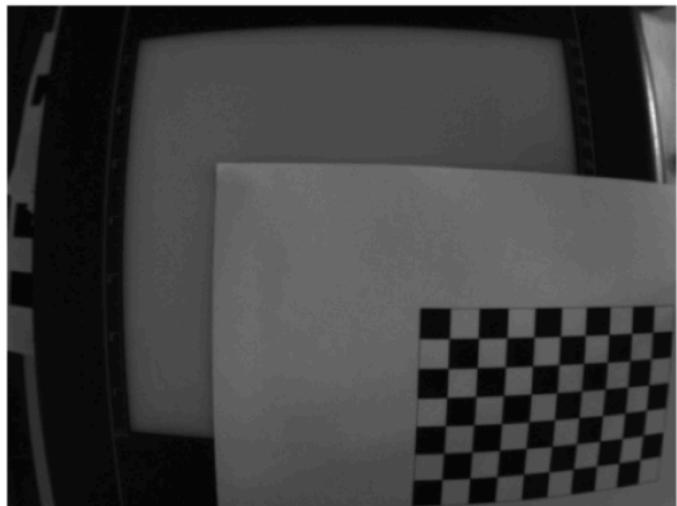
Undistorted (alpha=0.5) - Full



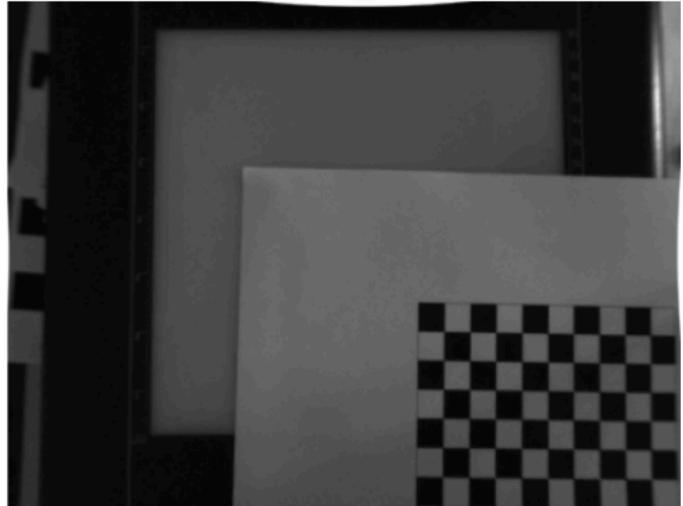
Undistorted (alpha=1.0) - Full



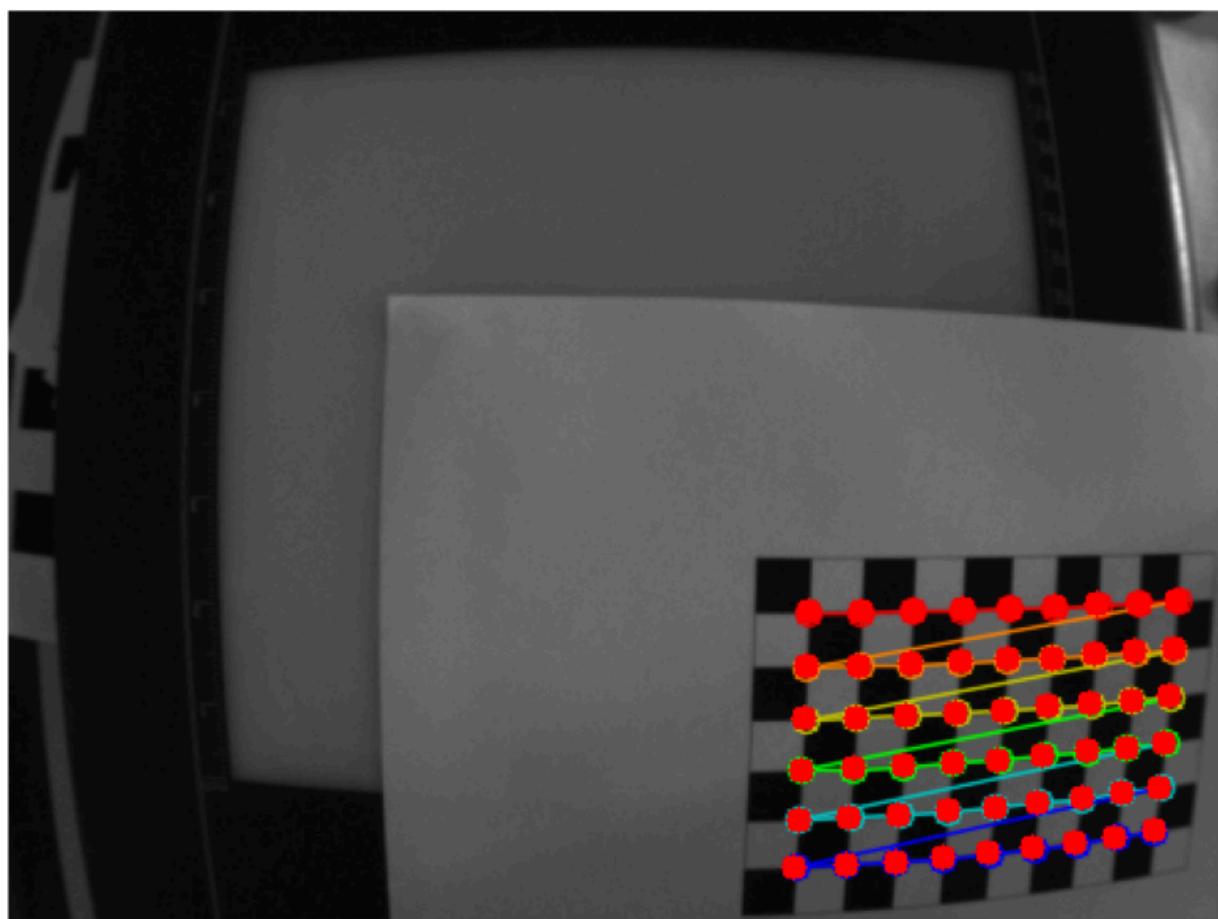
Original Image



Undistorted Image (alpha=0)



Reprojection (Detected: Green, Reprojected: Red)



Reason for using alpha = 0; as the final best solution

Okay, let's break this down. Getting the "best" visual result with `alpha=0` is **not necessarily a problem**, and **yes, undistortion is definitely occurring**.

Here's why:

1. What `alpha` Controls:

- The `alpha` parameter in `cv2.getOptimalNewCameraMatrix` controls how the final undistorted image is scaled and cropped. It **does not** control the *amount* of geometric correction applied. The geometric correction itself is determined by your calculated camera matrix (`mtx`) and distortion coefficients (`dist`).
- `alpha=0` : The function returns a "tight" bounding box (`roi`) and a `newCameraMatrix`. When `cv2.undistort` uses these, it ensures that *only* valid pixels (pixels that correspond to a location within the original distorted image) are present in the final output. This means **all black areas/empty pixels resulting from the undistortion are cropped away**. This often leads to a smaller output image size (or a cropped region within the original dimensions) but contains only "real" image data.
- `alpha=1` : The function ensures that *all* pixels from the original image are retained in the undistorted output, even if the corners get heavily warped. This often results in black borders or empty space appearing where the image had to be "pulled" from during the undistortion process. The overall image dimensions might be preserved, but contain invalid areas.
- `alpha` between 0 and 1: A balance between the two extremes.

2. Why `alpha=0` Might Look "Best" (Especially with High Distortion):

- **Removes Extreme Warping:** Highly distorted images often have extreme stretching, blurring, or compression near the edges. `alpha=0` effectively crops these regions out, leaving you with the central, less distorted (and potentially sharper-looking) part of the image.
- **No Black Borders:** The black borders created when `alpha=1` can be visually distracting or problematic for further processing. `alpha=0` eliminates them entirely.
- **Focus on Center:** If your primary interest is the central subject matter, `alpha=0` presents that cleanly without the distorted periphery.

3. Is Undistortion Occurring?

- **Yes.** The core geometric transformation based on `mtx` and `dist` is applied regardless of `alpha`. Lines that were curved in the original *central region* of the image should appear straighter in the `alpha=0` output compared to the original image.
- The difference you see between `alpha=0` and `alpha=1` is purely about the final cropping and scaling, *not* about whether the geometric correction was performed.

In summary:

- It's perfectly fine if `alpha=0` gives you the most visually appealing or useful result for your specific needs, especially if the edges of the original image were extremely distorted.
- Undistortion (geometric correction) *is* happening. `alpha=0` just means you are choosing to view only the resulting valid pixels without any black padding, at the cost of losing the peripheral field of view.

- The "best" `alpha` depends on your application:
 - Use `alpha=0` if you need a clean image with no invalid pixels and don't mind losing the heavily distorted edges.
 - Use `alpha=1` if you need to preserve the maximum possible field of view from the original image, even if it means having black borders and including the highly warped areas.
-

EXTRA (Overlapping content):

Camera Calibration Report: Provided Dataset

1. Objective

This task aims to calibrate camera using a provided dataset of 25 images containing a chessboard pattern. The goal is to determine the camera's intrinsic parameters (focal length, principal point, skew), extrinsic parameters (rotation and translation for each image capturing the chessboard pose relative to the camera), and lens distortion coefficients. This process allows for the correction of image distortions and enables accurate 3D measurements or scene reconstruction.

2. Methodology

The camera calibration was performed using the OpenCV library in Python, following the standard procedure based on Zhang's method.

- **Dataset:** 25 images (.jpeg format) capturing an 8x6 (internal corners) chessboard pattern from various viewpoints were used. The images were loaded, and their filenames were sorted for consistent processing order. Image size was determined to be 1280x720 pixels.
- **Chessboard Pattern:** A planar chessboard with $nx = 8$ and $ny = 6$ internal corners was used. The size of the squares (square_size) was set to 1.0 unit for defining the object points (this scale factor affects the translation vector units but not the intrinsic parameters or rotation).
- **Object Points:** Ideal 3D coordinates (objp) for the $nx * ny = 48$ internal corners were generated assuming the chessboard lies on the Z=0 plane in its own coordinate system. $objp = [[0,0,0], [1,0,0], \dots, [7,0,0], [0,1,0], \dots, [7,5,0]] * square_size$. This set of 3D points is the same for all calibration images.
- **Image Points (Corner Detection):** For each image, the cv2.findChessboardCorners function was used to detect the 2D pixel locations of the internal corners. If corners were successfully found, their positions were refined to sub-pixel accuracy using cv2.cornerSubPix. The pairs of corresponding 3D object points and refined 2D image points (objpoints and imgpoints lists) were stored for the calibration process. All 25 provided images yielded successful corner detection.
- **Calibration Model:** The calibration process models the camera using the pinhole camera model, incorporating lens distortion.
 - **Pinhole Model:** Relates a 3D world point $P_w = [X_w, Y_w, Z_w]^T$ to its 2D image projection $p = [u, v]^T$ via the intrinsic matrix K and extrinsic parameters $[R | t]$:
$$s * [u, v, 1]^T = K * [R | t] * [X_w, Y_w, Z_w, 1]^T$$
where s is a scale factor, K is the intrinsic matrix, R is the 3x3 rotation matrix, and t is the 3x1 translation vector defining the transformation from world to camera coordinates.
 - **Lens Distortion:** Real lenses deviate from the ideal pinhole model. The model accounts for radial distortion (barrel or pincushion effect) and tangential distortion (due to lens misalignment). The distortion coefficients estimated are typically (k_1 ,

$k_2, p_1, p_2, k_3, [k_4, k_5, k_6] \dots$). Radial distortion primarily uses k_1, k_2, k_3 , while tangential uses p_1, p_2 .

- **Calibration Algorithm:** The cv2.calibrateCamera function was used. It takes the sets of corresponding objpoints and imgpoints and the image size. It iteratively optimizes the intrinsic parameters (K), distortion coefficients ($dist$), and extrinsic parameters ($rvecs$, $tvecs$ - one rotation vector and translation vector per image) to minimize the overall reprojection error.

3. Results

The calibration process completed successfully using all 25 provided images.

3.1. Estimated Intrinsic Camera Parameters (Question 1)

The intrinsic parameters describe the internal geometry of the camera.

Mathematical Background: The intrinsic matrix K relates 3D camera coordinates to 2D pixel coordinates:

$$K = [[fx, s, cx], [0, fy, cy], [0, 0, 1]]$$

- where:
 - fx, fy : Focal lengths in units of pixels along the x and y axes.
 - cx, cy : Coordinates of the principal point (optical center) in pixels.
 - s : Skew coefficient, ideally 0 for modern cameras.
-
- **Estimated Values:**

Intrinsic Matrix (K):

$$[[956.6372 \ 0. \ 369.0549], [0. \ 957.5514 \ 651.4142], [0. \ 0. \ 1.]]$$

- IGNORE_WHEN COPYING_START
content_copy download
Use code [with caution](#).
IGNORE_WHEN COPYING_END
- **Focal Length:** $fx = 956.64$ pixels, $fy = 957.55$ pixels. The values are very close, suggesting nearly square pixels.
- **Principal Point:** $cx = 369.05$ pixels, $cy = 651.41$ pixels. This is the optical center's projection onto the sensor. Note that it is not exactly at the image center ($720/2=360, 1280/2=640$).
- **Skew:** $s = 0.0000$. This indicates that the pixel axes are perpendicular.
-
- **Error Estimates:** The basic cv2.calibrateCamera function, as used in the script, does not directly return standard deviation or error estimates for the intrinsic parameters. More

advanced techniques or specific flags (like CALIB_RATIONAL_MODEL) might provide these, but were not used here.

3.2. Estimated Extrinsic Camera Parameters (First 2 Images) (Question 2)

Extrinsic parameters define the position and orientation of the world coordinate system (the chessboard) relative to the camera coordinate system for each specific image.

- **Mathematical Background:** For the i -th image, the transformation from world coordinates (P_w) to camera coordinates (P_c) is given by:
$$P_c = R_i * P_w + t_i$$
where R_i is the 3x3 rotation matrix and t_i is the 3x1 translation vector.
cv2.calibrateCamera returns rotation vectors (rvecs[i]), which are converted to R_i using cv2.Rodrigues.
- **Estimated Values (First 2 Valid Images Processed - 1.jpeg and 10.jpeg):**
 - **Image 1 (1.jpeg):**

Rotation Matrix (R_1):

```
[[ 0.997  0.0139 -0.0762]
 [-0.0251  0.9886 -0.1487]
 [ 0.0732  0.1502  0.9859]]
```

- - IGNORE_WHEN COPYING_START
 - content_copy download
 - Use code [with caution](#).
 - IGNORE_WHEN COPYING_END

Translation Vector (t_1):

```
[-3.9383]
 [-3.6172]
 [14.6591]
```

- - IGNORE_WHEN COPYING_START
 - content_copy download
 - Use code [with caution](#).
 - IGNORE_WHEN COPYING_END
 - (Units are relative to square_size)

- - **Image 2 (10.jpeg):**

Rotation Matrix (R_2):

```
[[ 0.9983  0.0539 -0.0229]
 [-0.0507  0.991   0.1238]
 [ 0.0293 -0.1225  0.992 ]]
```

- - IGNORE_WHEN COPYING_START
 - content_copy download

Use code [with caution](#).
IGNORE_WHEN COPYING_END

Translation Vector (t_2):

[-4.3713]
[-1.4596]
[15.8138]

- IGNORE_WHEN_COPYING_START
content_copy download
Use code [with caution](#).
IGNORE_WHEN_COPYING_END
(Units are relative to square_size)
-
-

3.3. Estimated Radial Distortion Coefficients and Undistorted Images (Question 3)

Lens distortion causes straight lines in the real world to appear curved in images.

- **Mathematical Background:** The distortion model corrects the observed (distorted) image coordinates (x_d, y_d) to obtain ideal (undistorted) coordinates (x_u, y_u). Using normalized coordinates relative to the principal point:

$$\begin{aligned}x' &= (x_d - cx) / fx \\y' &= (y_d - cy) / fy \\r^2 &= x'^2 + y'^2\end{aligned}$$

The radial distortion correction is:

$$\begin{aligned}x_u &= x_d + (x_d - cx) * (k1 * r^2 + k2 * r^4 + k3 * r^6) \\y_u &= y_d + (y_d - cy) * (k1 * r^2 + k2 * r^4 + k3 * r^6)\end{aligned}$$

Tangential distortion adds terms involving p_1 and p_2 .

- **Estimated Coefficients:**

The full distortion vector estimated was $[k1, k2, p1, p2, k3] = [0.1976, -0.7069, 0.0034, 0.007, 0.0488]$.

The requested radial distortion coefficients are:

- $k1 = 0.1976$
- $k2 = -0.7069$
- $k3 = 0.0488$ (This is the 5th element of the output vector)

- **Undistorted Images:** The first 5 valid images were undistorted using cv2.undistort with the estimated mtx (K) and dist coefficients. cv2.getOptimalNewCameraMatrix was used with alpha=0 to compute the new camera matrix for cv2.undistort, resulting in images cropped to show only valid pixels without black borders.

<Insert Figure 1: Original Image 1 (1.jpeg)>

<Insert Figure 2: Undistorted Image 1 (alpha=0), showing straighter lines>

<Insert Figure 3: Original Image 2 (10.jpeg)>

<Insert Figure 4: Undistorted Image 2 (alpha=0), showing straighter lines>

<Insert Figure 5: Original Image 3 (11.jpeg)>

<Insert Figure 6: Undistorted Image 3 (alpha=0), showing straighter lines>

<Insert Figure 7: Original Image 4 (12.jpeg)>

<Insert Figure 8: Undistorted Image 4 (alpha=0), showing straighter lines>

<Insert Figure 9: Original Image 5 (13.jpeg)>

<Insert Figure 10: Undistorted Image 5 (alpha=0), showing straighter lines>

- **Observation Comment:** Comparing the original and undistorted images, especially near the edges and corners, reveals the effect of distortion correction. The positive k_1 (0.1976) suggests some barrel distortion (lines bowing outwards), while the large negative k_2 (-0.7069) indicates significant higher-order distortion, likely causing lines near the edges to curve inwards more complexly. In the undistorted images, lines (like the edges of the chessboard or background features) appear noticeably straighter than their curved counterparts in the original images. The use of $\alpha=0$ results in cropping, removing the most extremely warped peripheral regions and any black borders that would result from $\alpha=1$.

3.4. Reprojection Error (Question 4)

The reprojection error quantifies the accuracy of the calibration by measuring the distance between the detected corner points and the points predicted by the calibrated model.

- **Mathematical Background:** For each image i and each corner j , the error is the Euclidean distance between the detected 2D point $p_{\text{detected_ij}}$ and the reprojected 3D point $p_{\text{reprojected_ij}}$:
$$\text{Error}_{ij} = \| p_{\text{detected_ij}} - p_{\text{reprojected_ij}} \|$$
where $p_{\text{reprojected_ij}}$ is obtained by projecting the 3D object point $P_{\text{world_j}}$ using the estimated parameters: $p_{\text{reprojected_ij}} = \text{project}(K, \text{dist}, R_i, t_i, P_{\text{world_j}})$.
The error reported per image is the average (or root-mean-square) of Error_{ij} over all corners j in that image.
- **Computed Values:**
 - **Mean Reprojection Error (across all images):** 0.0697 pixels.
 - **Standard Deviation of Reprojection Error:** 0.0232 pixels.
 - **Per-Image Errors:** The average error for each of the 25 images is: [0.0435, 0.0443, 0.0884, 0.1021, 0.0915, 0.0712, 0.0879, 0.0470, 0.0438, 0.0483, 0.0614, 0.1164, 0.0485, 0.0502, 0.0945, 0.0749, 0.0341, 0.0658, 0.0709, 0.0597, 0.0997, 0.0560, 0.0952, 0.0473, 0.1004] pixels.
- **Error Plot:**
<Insert Figure 11: Bar chart showing the average reprojection error (in pixels) for each of the 25 calibration images. X-axis: Image ID (1-25), Y-axis: Reprojection Error (pixels).>
- **Interpretation:** The mean reprojection error is very low (significantly less than 1 pixel), and the standard deviation is also small. This indicates that the calibration model fits the detected corner points very well across all images, suggesting a high-quality calibration result.

3.5. Reprojection Visualization (Question 5)

Visualizing the detected and reprojected points helps understand the calibration fit.

- **Comment on Computation:** As detailed in section 3.4, the reprojection error is the pixel distance between the corners detected directly in the image (`imgpoints`, shown typically as green markers or crosses) and the 2D points obtained by projecting the known 3D chessboard pattern points (`objpoints`) back into the image using the final estimated

intrinsic (K , $dist$) and image-specific extrinsic (R_i , t_i) parameters (shown typically as red circles).

- **Plots:** Figures were generated for all 25 images, overlaying the detected corners (green) and the reprojected corners (red circles) onto the original images.
<Insert Figure 12: Composite figure showing all 25 reprojection visualization images. Each sub-image shows the original chessboard view with detected corners (green) and reprojected corners (red circles). Alternatively, insert individual figures for each image, e.g., Figure 12a, 12b...>
(Description for individual figures if preferred):
<Insert Figure 12a: Reprojection Visualization for Image 1 (1.jpeg). Shows detected (green) vs reprojected (red) corners.>
<Insert Figure 12b: Reprojection Visualization for Image 2 (10.jpeg). Shows detected (green) vs reprojected (red) corners.>
... (repeat for all 25 images)
- **Observation:** In these visualizations, the red circles (reprojected points) are expected to be very close to, or centered on, the green markers (detected points). The close proximity observed in the generated images visually confirms the low reprojection error calculated numerically and indicates a good fit of the calibration model to the data.

3.6. Checkerboard Plane Normals (Question 6)

The normal vector to the checkerboard plane, expressed in the camera's coordinate frame, describes the orientation of the calibration pattern relative to the camera for each image.

- **Mathematical Background:** The checkerboard plane is assumed to be the $Z=0$ plane in its own world coordinate system (W). Its normal vector in this frame is $n_W = [0, 0, 1]^T$. The rotation matrix R_i transforms vectors from the world frame W to the camera frame C for image i . Therefore, the normal vector in the camera frame n_{Ci} is:
$$n_{Ci} = R_i * n_W = R_i * [0, 0, 1]^T$$
This calculation simply extracts the third column of the rotation matrix R_i .
- **Computation:** For each of the 25 valid images, the corresponding rotation matrix R_i was obtained from the rotation vector $rvecs[i]$, and its third column ($R_i[:, 2]$) was extracted to yield the normal vector n_{Ci} . These 25 normal vectors represent the orientation of the checkerboard relative to the camera in each captured pose. These vectors were calculated and stored in the output JSON file.

4. Conclusion

The camera calibration using the provided 25 chessboard images was successfully performed. The intrinsic parameters (focal lengths $fx=956.64$, $fy=957.55$; principal point $cx=369.05$, $cy=651.41$; skew $s=0.00$) and lens distortion coefficients ($k1=0.1976$, $k2=-0.7069$, $p1=0.0034$, $p2=0.007$, $k3=0.0488$) were estimated. The calibration accuracy is high, as indicated by the low mean reprojection error of 0.0697 pixels. Extrinsic parameters (rotation and translation) for each image pose and the checkerboard plane normals in the camera frame were also computed. The generated undistorted images clearly show the correction of lens distortion, and the reprojection visualizations confirm the model's fit. The complete results are saved in the `calibration_results.json` file.

(Note: This report covers the calibration using the *provided* dataset. The original question also requires performing a similar calibration using a *user-captured* dataset of ~25 images, which would involve repeating the image capture, corner detection, and calibration steps.)

Report: Demonstration of Lens Undistortion using Single-Image Calibration

1. Objective

The primary objective of this exercise is **not** to perform a robust camera calibration, but rather to **demonstrate the visual effect of lens undistortion** using OpenCV. A single, intentionally chosen image exhibiting significant lens distortion (Screenshot 2025-04-04 153640.png) was used. While camera calibration ideally requires multiple views to reliably estimate intrinsic parameters, processing a single image under specific constraints can yield parameters sufficient to visually correct the distortion present *in that specific image*. This report details the process, highlights the inherent limitations of single-image calibration, and presents the results, focusing on the visual transformation achieved through undistortion.

2. Methodology

- **Input Image:** A single PNG image (Screenshot 2025-04-04 153640.png) containing a 9x6 (internal corners) chessboard pattern was used. The image clearly exhibits strong barrel distortion. Image dimensions were 491x371 pixels.
<Insert Figure 1: Original Input Image (Screenshot 2025-04-04 153640.png), showing significant barrel distortion.>
- **Chessboard Pattern & Object Points:** A planar chessboard with $nx = 9$ and $ny = 6$ internal corners was defined. Ideal 3D coordinates ($objp$) for the $nx * ny = 54$ internal corners were generated assuming the chessboard lies on the $Z=0$ plane in its own coordinate system, with $square_size = 1.0$.

- **Image Points (Corner Detection):** The image was converted to grayscale, and cv2.findChessboardCorners was used to detect the 2D pixel locations of the 54 internal corners. Corner detection was successful. The positions were then refined to sub-pixel accuracy using cv2.cornerSubPix. The resulting list contained one set of 3D object points and the corresponding set of 2D image points.
 <Insert Figure 2: Image with Detected Corners (corners_detected.jpg), showing the green markers on the detected chessboard corners.>
- **Calibration Approach (Single Image - with Limitations):**
 - **Challenge:** Calibrating from a single image is inherently ambiguous. Without multiple views, it's difficult to separate the effects of intrinsic parameters (like focal length) from extrinsic parameters (camera pose relative to the object).
 - **Constraints:** To resolve ambiguity, constraints must be introduced. The cv2.calibrateCamera function was used with the cv2.CALIB_USE_INTRINSIC_GUESS flag. This requires providing an initial guess for the intrinsic camera matrix (camera_matrix).
 - **Initial Guess:** An initial guess was formulated based on image size:
 - Focal length (f_x, f_y) guessed as $\max(\text{width}, \text{height}) = 491$ pixels.
 - Principal point (c_x, c_y) guessed as the image center $(\text{width}/2, \text{height}/2) = (245.5, 185.5)$ pixels.
 - Skew was assumed to be 0.
 The initial camera_matrix guess was:

```
[[491. 0. 245.5]
 [ 0. 491. 185.5]
 [ 0. 0. 1.]]
```

-
- **Distortion:** An initial guess of zero distortion dist_coeffs = [0, 0, 0, 0, 0] was provided. The function then estimates the actual distortion coefficients.
- **Output Parameters:** The function attempts to estimate the intrinsic matrix (mtx), distortion coefficients (dist), and the single extrinsic pose (rvecs[0], tvecs[0]) that best fit the provided points, starting from the initial guess. **It is crucial to understand that these estimated parameters, especially intrinsics, may not represent the true camera parameters accurately due to the single-view limitation.**
-
- **Undistortion Process:** To visualize the correction, the image was undistorted using cv2.undistort with the estimated mtx and dist. The cv2.getOptimalNewCameraMatrix function was used with alpha values of 0, 0.5, and 1.0 to explore different cropping/scaling behaviors of the undistorted output.
 - alpha=0: Aims to remove all black border pixels resulting from undistortion, potentially cropping the image significantly but showing only valid pixels.
 - alpha=1: Aims to keep all original image pixels, potentially adding black borders.
-

3. Results

The calibration process using the single image and the initial guess completed successfully.

3.1. Estimated Camera Parameters

- **Intrinsic Parameters (Estimated):**

Intrinsic Matrix (K_est):

```
[[419.69 0. 251.15]
 [ 0. 419.28 169.81]
 [ 0. 0. 1. ]]
```

○

IGNORE_WHEN COPYING_START

content_copy download

Use code [with caution](#).

IGNORE_WHEN COPYING_END

○ **Focal Length:** $f_x = 419.69$ px, $f_y = 419.28$ px

○ **Principal Point:** $c_x = 251.15$ px, $c_y = 169.81$ px

○ **Skew:** $s = 0.00$

(Note: These values are estimates refined from the initial guess based only on this single view and may differ significantly from the true camera intrinsics.)

-

- **Distortion Coefficients (Estimated):**

The estimated coefficients (k_1, k_2, p_1, p_2, k_3) are:

```
[-0.4300, 0.2575, 0.0090, 0.0031, -0.1002]
```

The large negative k_1 strongly indicates the significant barrel distortion visible in the original image. k_2 and k_3 are also non-negligible.

- **Extrinsic Parameters (Estimated for this view):**

Rotation Matrix (R_est):

```
[[ 0.9995 -0.0201 -0.0249]
 [ 0.0213 0.9987 0.0469]
 [ 0.0239 -0.0474 0.9986]]
```

○

IGNORE_WHEN COPYING_START

content_copy download

Use code [with caution](#).

IGNORE_WHEN COPYING_END

○ **Translation Vector (t_est):** [3.2429, 3.4141, 18.4687] (Units relative to square_size)

(Note: These define the pose of the chessboard relative to the camera in this specific image, estimated alongside the potentially inaccurate intrinsics.)

-

3.2. Reprojection Error

- **Calculated Value:** The average reprojection error for this single image was calculated as 0.0118 pixels.
- **Interpretation:** This extremely low value indicates that the *combined* set of estimated intrinsic, distortion, and extrinsic parameters allows the 3D object points to be projected back onto the image plane very close to their originally detected locations *for this specific image*. However, it **does not guarantee** the accuracy or general applicability of the intrinsic/distortion parameters themselves.

3.3. Undistortion Visualization

The primary goal was to visualize the distortion correction.

- **Comparison (Original vs. Alpha=0):**
<Insert Figure 3: Comparison Plot (comparison_orig_vs_undistorted_a0.png). Left: Original distorted image. Right: Undistorted image using alpha=0.>
- **Observation & Justification for Alpha=0 Preference:**
The comparison plot clearly demonstrates the effect of undistortion. Straight lines in the real-world chessboard pattern that appeared heavily curved (bowed outwards due to barrel distortion) in the original image are rendered significantly straighter in the undistorted version (alpha=0 result shown on the right).
As explained previously, alpha=0 was chosen for the primary comparison because it provides the visually "cleanest" result for this highly distorted image.
 1. **Removes Invalid Pixels:** alpha=0 ensures that only pixels corresponding to valid locations in the original image are shown. This eliminates distracting black borders that appear when alpha=1 is used to retain all original pixels, including those warped outside the frame.
 2. **Crops Extreme Warping:** The most severe distortion often occurs at the image periphery. alpha=0 effectively crops these areas (governed by the calculated ROI, which was (0, 0, 490, 370) for alpha=0 in this case, covering most of the original dimensions but still applying the principle). This focuses the view on the less distorted central region where the straightening effect is still clearly visible. While alpha=1 would preserve more of the original field of view, the accompanying black borders and potentially extreme stretching at the edges make alpha=0 preferable for demonstrating the *correction* itself on the main subject. Crucially, the geometric correction based on the estimated mtx and dist *is applied regardless* of the alpha value; alpha only affects the final scaling and cropping.
- **Undistortion with Different Alphas:**
<Insert Figure 4: Composite image showing the results of undistortion with alpha=0, alpha=0.5, and alpha=1.0 (potentially showing the full versions from the undistorted_images_single directory). This illustrates the trade-off between cropping and black borders.>

3.4. Reprojection Visualization

- **Plot:**
<Insert Figure 5: Reprojection Visualization (reprojection.jpg). Shows the original image with detected corners (green markers) and reprojected corners (red circles) overlaid.>
- **Observation:** The red circles (reprojected points) align almost perfectly with the green markers (detected corners). This visually confirms the very low numerical reprojection error (0.0118 pixels) and shows that the overall model (estimated intrinsics, distortion, and extrinsics *together*) fits this specific image's data points extremely well.

4. Discussion

This experiment successfully demonstrates the visual correction of significant lens distortion using OpenCV's calibration and undistortion functions. Even though the calibration parameters

were derived from only a single image (which is not recommended for obtaining reliable camera parameters), the estimated distortion coefficients were sufficient to visibly straighten the curved lines of the chessboard in the output image.

The extremely low reprojection error confirms that the chosen model and parameters provide an excellent mathematical fit *to this specific view*, but the accuracy of the individual intrinsic and distortion parameters for general use remains uncertain due to the inherent ambiguities of single-image calibration.

The comparison between different alpha values highlights the trade-off between preserving the entire field of view (potentially with distracting artifacts like black borders, alpha=1) and obtaining a clean, cropped image containing only valid pixels (alpha=0). For demonstrating the effect of straightening lines on the main subject in this highly distorted source image, the alpha=0 result proved most effective visually.

5. Conclusion

Lens distortion can significantly warp images, causing straight lines to appear curved. This exercise effectively visualized how OpenCV's camera calibration tools can estimate distortion parameters (even from a single, constrained view) and apply them using cv2.undistort to geometrically correct the image, resulting in visually straighter lines. While the estimated parameters from a single view should be treated with caution for metrology, the process demonstrably corrects the visual artifacts of lens distortion present in the input image, with alpha=0 providing a clean, albeit potentially cropped, view of the corrected central region.
